

AI THAT EXECUTES.

22 Frontier AI Model Guide & 10,000 Monthly Report Costing

Complete Rate Cards, Model Superpowers & Realistic Managed Costing for Arvocat Asset Managers

EXECUTIVE BRIEF FOR BOARD & TECHNICAL LEADERSHIP

ARVOCAP LEADERSHIP: Monicah Mwaniki (Co-Founder & CEO), John Ngure, Arnold Oduma (Tech Lead), Simar Juttla (Tech Lead)

C3A LABS TEAM: Satyam Singh Rajput (Product & Systems Lead), Pratyush Malviya (Sales Manager)

EXECUTIVE SUMMARY: Establishing report complexity tiers, evaluating 22 frontier AI models side-by-side (External Direct vs SARVAX Managed), and proving how Gemini OCR Ingestion + DeepSeek Text + Kimi K3 Brain delivers 10,000 monthly reports at \$950/mo. Delivery via Email & WhatsApp API.

NEXT NARRATIVE STEP: To scale AUM from KSh 11.02B to 20B+, we face a monthly client communication bottleneck.

Sending 10,000 Personalised Client Reports Every Month

10,000 Clients

Complete Investor Coverage

Every investor across Money Market, Almasi Fixed Income, and Equity funds receives a clear monthly update.

< 2 Minutes

Multi-Page PDF Generation

Replaces 48-hour manual analyst delays. Completes full PDF OCR, reasoning, and rendering in under 2 minutes per report.

Fixed Budget

Zero Surprise Pricing

Predictable \$950/mo flat rate. Fully managed end-to-end report generation, quality assurance, and automated multi-channel delivery.

THE COMMERCIAL STEP: ESTABLISHING REPORT BENCHMARKS BEFORE MODEL SELECTION

1. Page Count Drives Token Volume: A 2-page summary uses ~15,000 tokens, while a 5-page review uses ~35,000 tokens.
2. Why We Anchor on Tier 2 (Standard Review): Without exact sample PDFs, we establish Tier 2 (3–5 Pages @ \$950/mo) as our baseline benchmark.
3. Direct Email & WhatsApp Dispatch: Gemini Flash provides native OCR vision for scanned statements, with automatic delivery via Email & WhatsApp.

** Benchmark Footnote: Assumes an average workload of 35k input tokens and 3.5k output tokens per report (350M input / 35M output tokens across 10,000 monthly reports).*

NEXT NARRATIVE STEP: Before evaluating AI models, we define the 3 Report Complexity Tiers.

3 Commercial Tiers for 10,000 Monthly Reports (Managed Package)

TIER 1: LITE BRIEF

1 – 2 Pages

Short Monthly Summary

- Powered by Gemini 3.5 Flash-Lite (Fast OCR).
- Brief NAV & balance update statement.
- Volume: 15k input / 1.5k output per report (150M in / 15M out for 10k reports).
- Cost per Report: \$0.035 (~4.5 KSh / ₹3.38)
- 10k Monthly Total: \$350 / month
- Local Currency: ~45,100 KSh/mo (₹33.8k)

TIER 2: STANDARD REVIEW (RECOMMENDED)

3 – 5 Pages

Detailed Portfolio Breakdown

- Gemini OCR + DeepSeek Text + Kimi Brain.
- Full holding list, yield notes & market updates.
- Volume: 35k input / 3.5k output per report (350M in / 35M out for 10k reports).
- Cost per Report: \$0.095 (~12.2 KSh / ₹9.17)
- 10k Monthly Total: \$950 / month
- Local Currency: ~122,500 KSh/mo (₹91,700)

TIER 3: INSTITUTIONAL

8 – 12 Pages

Deep Institutional Analysis

- Gemini OCR + DeepSeek Text + Kimi K3 Brain.
- Comprehensive asset allocation & tax breakdown.
- Volume: 75k input / 8.0k output per report (750M in / 80M out for 10k reports).
- Cost per Report: \$0.195 (~25.1 KSh / ₹18.80)
- 10k Monthly Total: \$1,950 / month
- Local Currency: ~251,500 KSh/mo (₹1.88 Lakhs)

NEXT NARRATIVE STEP: Let's inspect the visual structure of a Standard 3-5 Page Client Report.

Anatomy of a Standard 3-5 Page Client Portfolio Report

PAGE 1 (OCR)

Portfolio Overview

- Ingested via Gemini Flash OCR
- Investor Name & Account ID
- Total Portfolio Value (KSh)
- Month-on-Month Growth (%)

PAGE 2 (OCR)

Asset Holdings

- Scanned Statement Parsing
- Almasi Fixed Income Balance
- Money Market Fund Balance
- Asset Class Pie Chart Vision

PAGE 3 (TEXT)

Yield & Returns

- DeepSeek Heavy Text Lifting
- Interest Yield Analysis
- Benchmark Comparison
- Distribution History

PAGE 4 (BRAIN)

Market Commentary

- Kimi K3 Financial Logic
- Kenyan Fixed Income Trends
- Central Bank Rate Impacts
- Risk & Outlook Brief

PAGE 5 (CONTROL)

Advisor Actions

- Rebalancing Suggestions
- Recommended Next Steps
- Assigned IFA Contact Details
- Click-to-Approve Gate

NEXT NARRATIVE STEP: Now let's compare 22 AI models side-by-side: Direct Uncached vs. SARVAX Managed Package.

High-Efficiency Context Readers (External Direct vs. SARVAX Package)

Model Name	Role / Superpower	Rates (In/Out per 1M)	External Direct 10k Cost	SARVAX Managed Package Price	SARVAX Savings
DeepSeek V4 Pro	★ Heavy Text Lifter (\$0.0036 cached)	\$0.435 / \$0.870	\$182.70 (~23.6k KSh)	\$150.00 (~19.3k KSh)	17.9% ★ Text Lead
Gemini 3.5 Flash-Lite	★ Primary OCR Reader (362 tps)	\$0.300 / \$2.500	\$192.50 (~24.8k KSh)	\$160.00 (~20.6k KSh)	16.9% ★ OCR Lead
Qwen 3.7 Plus	Base Input Rate (\$0.40) Table Extraction	\$0.400 / \$1.600	\$196.00 (~25.3k KSh)	\$165.00 (~21.2k KSh)	15.8%
Cohere Command A	Enterprise RAG & Source Citations	\$0.500 / \$1.500	\$228.00 (~29.4k KSh)	\$190.00 (~24.5k KSh)	16.7%
Amazon Nova Pro	AWS Bedrock Enterprise VPC Security	\$0.800 / \$3.200	\$392.00 (~50.6k KSh)	\$310.00 (~40.0k KSh)	20.9%
Kimi K2.6	Long Context Under \$1.00 Input Rate	\$0.950 / \$4.000	\$472.50 (~61.0k KSh)	\$370.00 (~47.7k KSh)	21.7%
GPT-5.6 Luna	Sub-Second Low Latency Chat Copilot	\$1.000 / \$6.000	\$560.00 (~72.2k KSh)	\$450.00 (~58.0k KSh)	19.6%

* Benchmark Volume Footnote: Calculated for Tier 2 Standard (35k input tokens + 3.5k output tokens per report = 350M input / 35M output tokens across 10,000 reports).

NEXT NARRATIVE STEP: Part 2 compares financial reasoning models and multimodal vision models.

Financial Reasoning & Vision Models (External Direct vs. SARVAX Package)

Model Name	Role / Superpower	Rates (In/Out per 1M)	External Direct 10k Cost	SARVAX Managed Package Price	SARVAX Savings
Kimi K3 (Moonshot)	★ #1 Global SOTA TAU Banking (0.3340)	\$3.000 / \$15.000	\$1,575.00 (~203.2k KSh)	\$950.00 (~122.5k KSh)	39.7% ★ Brain Lead
Gemini 3.5 Flash	Fixed 258 Tokens/Page Scanned PDF Vision	\$1.500 / \$9.000	\$840.00 (~108.4k KSh)	\$620.00 (~80.0k KSh)	26.2%
Gemini 3.6 Flash	1M Native Context & OCR Speed	\$1.500 / \$7.500	\$787.50 (~101.6k KSh)	\$550.00 (~70.9k KSh)	30.2%
Qwen 3.7 Max	SOTA Multi-Sheet Excel Logic (0.3120 TAU)	\$2.500 / \$7.500	\$1,137.50 (~146.7k KSh)	\$780.00 (~100.6k KSh)	31.4%
GLM-5.2	Multi-Turn Dialogue Context Retention	\$1.400 / \$4.400	\$644.00 (~83.1k KSh)	\$510.00 (~65.8k KSh)	20.8%
Grok 4.5 (xAI)	Real-Time Market Search Integration	\$2.000 / \$6.000	\$910.00 (~117.4k KSh)	\$690.00 (~89.0k KSh)	24.2%
Claude Sonnet 5	Executive Writing Tone for UHNW Briefs	\$2.000 / \$10.000	\$1,050.00 (~135.5k KSh)	\$750.00 (~96.7k KSh)	28.6%

* Benchmark Volume Footnote: Calculated for Tier 2 Standard (35k input tokens + 3.5k output tokens per report = 350M input / 35M output tokens across 10,000 reports).

NEXT NARRATIVE STEP: Part 3 compares heavy frontier flagship models.

Heavy Frontier & Enterprise Models (External Direct vs. SARVAX Package)

Model Name	Role / Superpower	Rates (In/Out per 1M)	External Direct 10k Cost	SARVAX Managed Package Price	SARVAX Savings
Gemini 3.1 Pro Preview	Deep Multimodal Code & Analysis	\$2.000 / \$12.000	\$1,120.00 (~144.5k KSh)	\$820.00 (~105.8k KSh)	26.8%
GPT-5.4 (OpenAI)	High Function-Calling Tool Reliability	\$2.500 / \$15.000	\$1,400.00 (~180.6k KSh)	\$980.00 (~126.4k KSh)	30.0%
GPT-5.6 Terra	High Concurrency Limits & Uptime SLA	\$2.500 / \$15.000	\$1,400.00 (~180.6k KSh)	\$980.00 (~126.4k KSh)	30.0%
Claude Sonnet 4.6	Flawless Document Formatting & Artifacts	\$3.000 / \$15.000	\$1,575.00 (~203.2k KSh)	\$980.00 (~126.4k KSh)	37.8%
Claude Opus 5	Ultimate Frontier Reasoning Flagship	\$5.000 / \$25.000	\$2,625.00 (~338.6k KSh)	\$1,750.00 (~225.8k KSh)	33.3%
GPT-5.5 (OpenAI)	Heavy Frontier Code & Math Logic	\$5.000 / \$30.000	\$2,800.00 (~361.2k KSh)	\$1,950.00 (~251.5k KSh)	30.4%
Claude Fable 5	Specialized Long-Form Narrative Research	\$10.000 / \$50.000	\$5,250.00 (~677.3k KSh)	\$3,500.00 (~451.5k KSh)	33.3%

* Benchmark Volume Footnote: Calculated for Tier 2 Standard (35k input tokens + 3.5k output tokens per report = 350M input / 35M output tokens across 10,000 reports).

NEXT NARRATIVE STEP: Rather than paying expensive standalone rates, SARVAX uses a Smart Multi-Model Cascade.

Smart Multi-Model Cascade: OCR Vision + Text Heavy Lifting

STEP 1: OCR VISION

Gemini Flash (OCR)

Scanned PDF Vision Reader

- Scans image PDFs, identity docs & charts.
- Fixed 258 tokens per page.
- Extracts structured text & numbers from scanned files.

STEP 2: HEAVY TEXT LIFTING

DeepSeek V4 Pro

Text Context Ingestion

- Ingests 50+ page CRM histories & fund disclosures.
- 99.17% Caching discount (\$0.0036/1M cached).
- High-volume context processing at pennies.

STEP 3: DEEP REASONING

Kimi K3 Brain

#1 SOTA Banking Reasoning

- #1 TAU Banking benchmark leader (0.3340).
- Writes clear, professional financial commentary.
- Explains yield & market shifts clearly.

STEP 4: EXACT MATH & CONTROL

Exact Math Oracle

100% Exact Math + Advisor Sign-Off

- Python Decimal code handles 100% of NAVs & returns (Zero AI math error).
- Human advisors review & approve drafts.
- Full audit trail kept for internal records.

NEXT NARRATIVE STEP: See how Smart Memory Caching slashes monthly API bills by another 65%.

- COST SAVINGS

Smart Caching & Platform Economics: Direct Vendor vs. SARVAX Benchmark

DIRECT VENDOR UNCACHED RATE

\$1,575 / Month

~203,200 KSh / mo • ₹1.52 Lakhs / mo

Standalone Direct API Pricing

$\$3.00/1M \text{ In } (\$0.105) + \$15.00/1M \text{ Out } (\$0.0525) = \$0.1575 / \text{report}.$

Zero Ingestion Optimization

Buying direct from vendor API charges full rate on every run.

10,000 Uncached Runs Total

Assumes 350M input / 35M output tokens across 10,000 reports.

SARVAX MANAGED TIER 2 PACKAGE

39.7% CLIENT SAVINGS

\$950 / Month

~122,500 KSh / mo • ₹91,700 / mo

All-Inclusive Managed SaaS

Fully managed report generation, quality checks, and investor distribution.

Kimi K3 SOTA Banking Engine

Includes #1 SOTA TAU Banking reasoning + Gemini OCR Vision reader.

Safe Flat-Rate Guarantee

Saves \$625/mo vs direct vendor. Capped flat rate guarantees complete budget predictability.

NEXT NARRATIVE STEP: How batch queueing ensures month-end reports dispatch reliably.

Automated Monthly Dispatch & Encrypted Data Isolation

01 AUTOMATED MONTH-END QUEUEING

Runs quietly in the background at month-end. System processes all 10,000 client reports in under 3.5 hours without requiring manual intervention from operations or IT.

02 PRIMARY EMAIL & WHATSAPP API DISPATCH

Dispatches PDF reports directly to investor email inboxes and WhatsApp API messages. Zero integration dependencies on mobile app development.

03 ENCRYPTED DATA PRIVACY & ZERO MODEL TRAINING

Strict Zero Data Retention agreements with LLM providers. Client financial data is encrypted in dedicated isolated storage and NEVER trained on by public AI models.

NEXT NARRATIVE STEP: The bottom-line financial ROI and recommended action items for today's call.

2,300+ Analyst Hours Recovered at <1% Cost & Pilot Plan

QUANTIFIED FINANCIAL ROI

2,340 Hours / Year

Recovered Across Wealth Advisors & Operations

- Manual Alternative: Hiring 3 middle-office analysts = ~\$45,000/yr (~5.8M KSh/yr).
- SARVAX Standard Tier 2 Cost: ~\$11,400 / year (~1.47M KSh/yr / \$950/mo).
- Direct Staffing Savings: >74.7% savings vs traditional analyst staffing.
- 100% Investor Coverage: Guarantees every single client receives a statement.

RECOMMENDED PACKAGE & NEXT STEPS FOR CALL

Tier 2: Standard Package

~\$950 / Mo (~122,500 KSh / ₹91,700/mo) for 10k Reports

1. Confirm Tier Selection: Select Tier 1 (\$350/mo), Tier 2 (\$950/mo), or Tier 3 (\$1,950/mo).
2. Share 1 Sample Report PDF: Allows C3A engineering to lock exact template bounds.
3. Run 100-Report Free Test Batch: Verify Gemini Flash OCR & Email/WhatsApp dispatch.
4. Kickoff Technical Setup: Connect BFF API endpoints with Arnold & Simar.

NEXT NARRATIVE STEP: Let's review C3aLabs platform overview and leadership contact details.



C3aLabs, Inc. × Arvocat Asset Managers

Autonomous Execution & Communication Engine for Wealth Advisory

SarvaX.ai

Agentic AI Platform for Regulated Wealth & Asset Management

Sai Casula — Founder & CEO

saicasula@ c3alabs.com · c3alabs.com · New York, NY

C3A Project Contacts: Satyam Singh Rajput (Product Lead) | Pratyush Malviya (Sales Manager)

"AI won't replace financial advisors. But advisors empowered by AI will serve 10x more investors with higher trust."