

The trusted decision engine for institutional wealth management. We help C3A Labs and financial leaders design multi-model topographies (24 models across 40 configurations), justify unit token economics, deploy deterministic control flows, enforce SEBI/FINRA/SOC2 governance, and continuously optimize production workloads.

PILLAR 1

DESIGN

Multi-Model Reader-Brain topologies and hybrid provider routing.

PILLAR 2

JUSTIFY

Exact INR token economics, 25-model rates, and Lakhs/Crores ROI oracles.

PILLAR 3

DEPLOY

Deterministic DAG workflows, advisor approval gates & CRM writebacks.

PILLAR 4

GOVERN

SEBI, FINRA 3110, SEC 204-2, SOC2 Type II, and INT4 vs FP8 safety gates.

PILLAR 5

OPTIMIZE

Continuous context normalization, prompt caching, and latency tuning.

Evidence:

Verdict: Conditional Hybrid Cascade — Evidence Gates Remain

A hybrid cascade remains a valid **architecture pattern**, not a verified model-selection decision. Twenty-one unique models have both provider-family documentation and at least one exact retained benchmark row; three need more proof. Financial savings, provider rate limits and production SLAs require exact SKU, contract and trace evidence before GTM approval.

Decision: architecture exploration may continue; procurement, financial commitment and regulated production routing remain conditional.

Evidence boundary: provider pages establish model-family documentation. Artificial Analysis metrics are from a retained API snapshot that could not be freshly refreshed because the API credential was unavailable. A benchmark row does not establish provider availability, SLA, compliance or customer access.

21DUAL-PROOF MODEL
CARDS**3**WEALTH-ADVISORY
JOURNEYS**12**DOMAIN AGENT
ROLES**3**PENDING
VERIFICATION

Executive Product Verdict: Wealth Advisory Architecture

A **Product and Enterprise Margins POV** evaluating the frontier landscape for long-running financial wealth workflows.

SARVAX WEALTH-ADVISORY ARCHITECTURE

12 agents · 3 advisory journeys

Models underneath the product—not the product itself

Wealth advisory work arrives through meetings, CRM history, email, WhatsApp, holdings, market data, KYC records, tasks and client documents. SARVAX owns the journey: intake, agent orchestration, memory, approvals, integrations and writeback. Foundation models perform bounded routing, context assembly, vision and reasoning roles underneath it.

Product recommendation: Design around advisor journeys—not PDF types. Keep deterministic calculations, policy enforcement, authorization and human approval outside the model chain.

Relationship intake

Meeting Assistant + CRM/MCP context



The Router

Gemini 3.5 Flash-Lite



The Reader

DeepSeek V4 Pro



Relationship agents

Pre/Post Meeting + Client Strategist



Advisor approval

Workflow 2.0 approval gate



System writeback

CRM, tasks, notifications, artifacts

Turns meetings and ongoing client communication into advisor-ready context, follow-ups and controlled system updates.

SELECTED SARVAX JOURNEY

Client Relationship Intelligence

ADVISOR OUTCOME

Every advisor starts prepared, leaves with structured actions and keeps the client record current.

SARVAX AGENTS

Pre-Meeting Analysis Agent

Post-Meeting Analysis Agent

Client Intelligence Strategist

PRODUCT SURFACES

Meeting Assistant

Agent Builder

Agent Memory

Workflow 2.0

MCP / integrations layer

Notifications and writeback

INPUTS ACROSS THE JOURNEY

Meeting Assistant transcripts

Email and WhatsApp context

CRM household history

Advisor notes and tasks

CONTROL BOUNDARY

Meeting intelligence drafts context and actions. The advisor approves material client communication and every financial recommendation.

WEALTH ADVISORY JOURNEY	ASSIGNED MODEL STACK	INPUT TOKENS	OUTPUT TOKENS	AVG INPUT COST (₹)	AVG OUTPUT COST (₹)	COST / RUN (₹)	MONTHLY COST (10K RUNS)
	Flash-Lite → DeepSeek V4 → Kimi K3	70,000	6,200	₹5.61	₹2.80	₹8.42	₹0.84 Lakhs / mo
	DeepSeek V4 → Kimi K3	135,000	11,500	₹20.53	₹9.15	₹29.69	₹2.97 Lakhs / mo
	Gemini Vision → DeepSeek V4	70,000	5,500	₹1.84	₹0.38	₹2.22	₹0.22 Lakhs / mo

Mathematical Note: Costs are calculated per execution step using high-precision Decimal arithmetic and official provider rate cards. Input tokens are cheaper (\$0.15–\$3.00/1M) due to parallel GPU prefill; output tokens are 3x–8x costlier (\$0.60–\$15.00/1M) due to sequential generation.

BOUNDED MODEL ROLES UNDERNEATH THESE JOURNEYS

CORE STAGE 2

Moonshot AI

Product role: Produces bounded portfolio, tax and compliance analysis drafts inside SARVAX agents. Deterministic calculators and advisor approval own the final decision.

Why: Its retained TAU Banking result is 0.3340, the highest value in this report’s retained 40-configuration comparison set. This is a shortlist signal, not proof of production accuracy.

Boundary: Use only where a wealth-advisory journey needs deeper financial reasoning. Meeting notes, routing and routine extraction should not pay this cost or inherit this latency.

0.3340

TAU Banking
Retained independent snapshot

33.1 tps

Speed
Retained independent snapshot

₹289.70

Input / 1M
Retained benchmark rate; exact SKU price
needs claim-level verification

Task testing required

Provider proof • Benchmark

DeepSeek

Product role: Normalizes high-volume text from meeting transcripts, email, WhatsApp exports, CRM notes, research and text-native statements before SARVAX agents reason over it.

Why: The official rate card verifies the exact SKU, low cache-hit pricing, JSON output, tool calls and a 1M context window—useful for large client-context assembly.

Boundary: It is a text-context worker, not the product. Image-origin evidence still needs the Vision Gate; extracted facts must retain source references and pass journey-specific validation.

0.2371–0.2577

TAU Banking range
Retained independent configurations

31.2–44.3

General score range
Retained independent configurations

₹42.01

Cache-miss input / 1M
Official DeepSeek rate card

Official SKU and price verified; task quality unproven

[Provider proof](#) • [Benchmark](#)

Google

Product role: Identifies the wealth-advisory journey, urgency, household context and required SARVAX agent/workflow before expensive processing begins.

Why: Google describes it as its fastest, most cost-effective 3.5 model for high-throughput execution. The retained snapshot reports 362.2 tps; its verified input rate is ₹28.97/1M.

Boundary: Use for routing, classification and low-risk extraction checks—not investment, tax or compliance conclusions. Low-confidence routing must fall back to review.

362.2 tps

Speed
Retained independent snapshot

₹28.97

Input / 1M
Verified official rate card

0.1649

TAU Banking
Retained independent snapshot

Conditional gate; task testing required

[Provider proof](#) • [Benchmark](#)

Google

Product role: Handles image-origin evidence—scanned KYC records, photographed forms, charts and statement pages—when a broader wealth-advisory journey actually needs it.

Why: Google’s model page describes Gemini 3.6 Flash as supporting agentic and multimodal tasks. The retained snapshot reports 243.9 tps.

Boundary: Multimodal identity does not prove financial-document OCR accuracy. Validate tables, decimal placement, signatures and low-quality scans; escalate low-confidence evidence.

243.9 tps Speed Retained independent snapshot	50.1 General score Retained independent snapshot	Exact SKU rate not established Price Retained benchmark rate; exact SKU price needs claim-level verification
--	---	---

Multimodal identity verified; OCR accuracy unproven

Provider proof • Benchmark

► Evidence boundaries and QA corrections

SECTION 2

SARVAX Wealth-Advisory Agent Operating Matrix

Verified production client accounts, assigned developers, active integrations, and empirical token unit economics sourced directly from official specification documents and Google Sheets.

9 Active Accounts · Source-Linked to Google Drive

CLIENT ACCOUNT	OWNER / DEVELOPER	STATUS	PRIMARY INTEGRATION	REAL WORKLOAD SCOPE	ASSIGNED MODEL	UNCACHED COST / RUN	80% CAPACITY COST RUN
1oU82tIF... (Doc).	Yash Bhorkar / Ria Choudhari	IN TESTING	Shared Email Inbox + Excel	50 daily lab reports, Micro vs Analytical PDF parsing, OOS color coding, 6-mo skip-test check, 8 AM summary email.	Gemini 3.5 Flash (Vision)	₹2.32 / report	₹0. rep
13UCF0P2... (Doc).	Sarang Kulkarni / Deviprasad	IN DEVELOPMENT	Periskope WhatsApp + Sheets	200 dedicated client WhatsApp groups, 3-hour refresh cycle, Aadhaar / PAN / Passport OCR document extraction.	Flash-Lite → DeepSeek V4	₹0.72 / check	₹0. che
1DGThpbn... (Doc).	Sarang Kulkarni / Ria Choudhari	COMPLETED	OneChat + Agent Builder	10k monthly client reports, investment briefs, and CMA Nairobi compliance.	DeepSeek V4 → Kimi K3	₹17.96 / brief	₹6. bri
1ApSkZww... (Sheet Row 19).	Shubhangi / Siddhi Satam	READY FOR DEMO	OneChat + Advisory Workflows	In-person meeting completed (Suresh & Rohit). Goal planning, asset allocation, and advisor task automation.	Flash-Lite → DeepSeek V4	₹1.11 / run	₹0. run
1ApSkZww... (Sheet Row 20).	Sarang / Ayanabha Misra	REQUIREMENT REVIEW	Workflows 2.0 + Zoho Projects	POC usecases defined by dev team; awaiting Sai's	Flash-Lite → DeepSeek V4	₹2.35 / run	₹0. run

CLIENT ACCOUNT	OWNER / DEVELOPER	STATUS	PRIMARY INTEGRATION	REAL WORKLOAD SCOPE	ASSIGNED MODEL	UNCACHED COST / RUN	80' CA' CO RU
				review. Wealth onboarding and rebalancing workflows.			
1ApSkZww... (Sheet Row 6).	Sarang / Ria Choudhari	BLOCKED (CLIENT FEEDBACK)	OneChat + Wealth Workflows	POC demo completed with Vinit Iyer on Jul 24. Awaiting client feedback and next steps.	DeepSeek V4 → Kimi K3	₹8.42 / run	₹2. run
1ApSkZww... (Sheet Row 13).	Shubhangi / Naman Jaiswal	POC IN PROGRESS	WealthSpectrum MCP + Morningstar	RM Client 360 View: Pre-meeting briefs, live portfolio analysis, client research, and post-meeting follow-ups.	DeepSeek V4 → Kimi K3	₹8.42 / run	₹2. run
1ApSkZww... (Sheet Row 22).	Pratyush / Deviprasad	TOOL R&D	Syntoniq, Investwell, Ace MF	Bengaluru meeting Jul 29 (Sai & Avantika). Multi-platform connectors for order placement and portfolio normalization.	Gemini 3.5 Flash	₹3.48 / run	₹1. run
1ApSkZww... (Sheet Row 23).	Sarang / Siddhi / Naman	TOOL R&D	Morningstar, Reuters, Bloomberg	Institutional family office market feed synthesis, risk reporting, and custodian statement ingestion.	Gemini 3.5 Flash	₹5.22 / report	₹1. rep

Primary Source Verification: Sourced directly from C3A Labs Google Drive specifications (`1oU82tIF...`, `13UCF0P2...`, `1DGThpbn...`) and Google Sheet `Usecase Tracker` (`1ApSkZww...`). Exchange basis: ECB FX ₹96.57 / \$1.

Sales Team Playbook · Non-Technical Primer

THE 3-LAYER PITCH

Productivity vs. Agentic vs. Platform AI

When pitching to a CEO or Head of Wealth, divide SARVAX into 3 simple layers:

Layer 1 (Productivity): What the advisor sees — audio transcription, WhatsApp OCR, auto-filling CRM tasks.

Layer 2 (Agentic): Bounded specialized agents (Pre-Meeting Agent, Portfolio Analyst) that ingest data and write drafts.

Layer 3 (Platform & Guardrails): Workflows 2.0 DAG engine enforces exact Decimal math, policy checks, and mandatory advisor sign-off.

TOKEN ECONOMICS ELI5

The "Reading vs. Writing" Analogy

Explain why input costs dominate output costs in financial advisory:

Reading (Input): The AI reads 50 pages of portfolio statements + 3 years of CRM history (75,000 tokens). Reading is cheap (₹14.50–₹29.00/1M).

Writing (Output): The AI writes a concise 2-page rebalancing summary (8,000 tokens). Writing is 5x costlier due to sequential GPU generation.

Key Takeaway: Reading volume is 10x higher, making input tokens the main driver of overall run costs.

COST SAVINGS CHEAT-SHEET

How Prompt Caching Slashes Client Invoices

Use these exact numbers during commercial pricing calls with prospects:

Uncached Standard Rate: Portfolio review = ₹29.69 / run (₹2.97 Lakhs/mo for 10k runs).

With 80% Prompt Caching: Cost drops to ₹10.97 / run (₹1.10 Lakhs/mo) — 63% client savings!

WhatsApp Automation (NS Wealth): Cost is just ₹0.24 / check (only ₹57.60/client/month for 24/7 coverage).

OBJECTION HANDLING PLAYBOOK

Answering CXO Security & Accuracy Questions

Two core objections every wealth manager asks during a demo:

"Will the AI hallucinate portfolio returns?" — *Answer:* No. In SARVAX, models never do math. Our Python Decimal engine calculates 100% exact return numbers.

"Can AI send emails without approval?" — *Answer:* No. Every outbound action hits an `await_approval` gate requiring human advisor click-to-confirm.

Standard-Token Cost Scenario

Compares standard token rates across frontier models in the repository catalog using exact Decimal arithmetic and the official ECB USD/INR reference rate. Workload volumes are editable planning assumptions—not production invoices or observed SARVAX usage.

Verified inputs

Catalog rate cards and ECB 2026-07-24 USD/INR cross-rate.

Scenario inputs

Input tokens, output tokens and monthly execution volume.

Excluded

Cache, batch, retries, tools, storage, platform fees, taxes and negotiated contracts.

Formula: $((\text{input tokens} \times \text{input USD/1M}) + (\text{output tokens} \times \text{output USD/1M})) \times \text{USD/INR}$. Model-specific context and pricing tiers are enforced before a result is shown.

INPUT TOKENS What the AI Reads

The prompt, background context, client files, and system instructions ingested by the model.

Real-Life Wealth Advisory Example (75,000 Tokens):

An advisor executes a portfolio review for an HNI client. SARVAX ingests:

45-minute Meeting Audio Transcript (~10,000 tokens)

3 Years of CRM Client Notes & IPS Guidelines (~15,000 tokens)

WealthSpectrum Portfolio & Holdings Data (~50,000 tokens)

Pricing Impact: Input tokens are cheaper (\$0.30–\$3.00 / 1M) because processing prompt text in parallel is GPU-efficient.

OUTPUT TOKENS What the AI Generates

The response text, drafted emails, recommendations, and structured JSON generated by the model.

Real-Life Wealth Advisory Example (8,000 Tokens):

SARVAX agents synthesize the findings and output:

Personalized Follow-Up Email & WhatsApp Note (~1,500 tokens)

Rebalancing Proposal & Asset Allocation Report (~4,500 tokens)

Structured CRM Task Records & Audit Logs (~2,000 tokens)

Pricing Impact: Output tokens are 3x–8x costlier (\$1.60–\$15.00 / 1M) because tokens are generated sequentially token-by-token.

Planning inputs

SCENARIO PRESET

Portfolio Review (75,000 in / 8,000 out)▼

MODEL A

Gemini 3.5 Flash-Lite (Google) — \$0.30 in / \$2.50 out per 1M [Verified]▼

MODEL B

DeepSeek V4 Pro (DeepSeek) — \$0.435 in / \$0.87 out per 1M [Verified]▼

Input tokens per run75,000

Output tokens per run8,000

Monthly execution volume10,000 runs

FX basis: ECB 2026-07-24 — ₹96.567636/USD. Latest business-day official cross-rate.

Calculated with ECB 2026-07-24 FX and verified public rate inputs. Gemini 3.5 Flash-Lite: Global standard inference; official Google rate card. DeepSeek V4 Pro: Standard cache-miss input and output; official DeepSeek rate card.

Calculated standard-token cost

Global standard inference; official Google rate card
PER RUN

₹4.10

MONTHLY

₹0.4104 lakh / month

PER 1,000 RUNS

₹4104.12

Standard cache-miss input and output; official DeepSeek rate card
PER RUN

₹3.82

MONTHLY

₹0.3823 lakh / month

PER 1,000 RUNS

₹3822.63

CALCULATED PRICE DIFFERENCE

DeepSeek V4 Pro is lower for these planning inputs. This is a rate-only difference—not measured savings.

ANNUALIZED DIFFERENCE

INTERPRETATION

SECTION 4

Wealth-Advisory Control-Flow Sandbox

Tests where product orchestration, deterministic gates and advisor approval must sit before a POC is approved. It demonstrates control order—not production execution quality.

Client Relationship Intelligence



Scenario progress

Ready

Step Costing & Telemetry

Executed: 0 in / 0 out (₹0.00) · Total Run: ₹8.42 · Monthly: ₹0.84 Lakhs/mo (10k runs)

DECISION UNDER TEST

Can SARVAX preserve evidence, ownership and human decision rights across the complete meeting journey?

From permitted client context to advisor-approved follow-up and controlled system writeback.

PRODUCT OWNER

Meeting Assistant + Workflow 2.0

DETERMINISTIC OWNER

Policy, authorization and write-integrity services

HUMAN DECISION RIGHT

Named advisor

EXIT GATE

Release only after consent, tenant isolation, source lineage, approval and writeback controls pass representative trace tests.

RELEASE EVIDENCE STILL REQUIRED

Connector execution, persistence, retry behavior and production reliability require representative backend traces.

Control-flow chain

Accurate Token & Costing Model — Verified Model Rates (ECB ₹96.57/\$1)

Pending Active Passed

1 Permitted client-context intake PRODUCT PENDING

Model: Gemini 3.5 Flash-Lite · ₹0.14 (5,000 in / 0 out)

Calendar, CRM and approved communication context are collected before preparation.

Required gate: Required: consent, tenant access, source timestamps and retention policy.

2 Pre-Meeting Analysis Agent AGENT PENDING

Model: DeepSeek V4 Pro · ₹0.73 (15,000 in / 1,200 out)

Drafts a source-linked preparation brief and explicit data gaps.

Required gate: Required: unsupported client facts blocked; every material claim links to retained evidence.

3 Meeting capture and artifact finalization PRODUCT

PENDING

Model: Gemini 3.5 Flash-Lite · ₹0.88 (18,000 in / 1,500 out)

Meeting Assistant creates the permitted transcript and meeting artifacts.

Required gate: Required: participant consent, speaker attribution review and artifact access policy.

4 Post-Meeting Analysis Agent AGENT

PENDING

Model: DeepSeek V4 Pro · ₹1.01 (20,000 in / 2,000 out)

Drafts facts, decisions, risks and action items from the retained artifact.

Required gate: Required: transcript facts remain separate from model interpretation.

5 Client Intelligence Strategist AGENT

PENDING

Model: Kimi K3 · ₹5.65 (12,000 in / 1,500 out)

Drafts next-best-action options for advisor review using TAU SOTA reasoning.

Required gate: Required: no autonomous product recommendation, promise or client communication.

6 Advisor decision gate HUMAN

PENDING

Model: Named Advisor Gate · Free / Deterministic

The named advisor accepts, edits or rejects facts and proposed actions.

Required gate: Required before any client communication or material CRM/task update.

7 Controlled Workflow 2.0 writeback TOOL

PENDING

Model: Workflow 2.0 API · Free / Deterministic


Only approved records are submitted to authorized connected systems.

Required gate: Release evidence required: idempotency key, immutable audit event, retry/dead-letter policy and compensating action.


Four Evidence-Backed Comparison Charts

Charts read the central runtime dataset. Each chart discloses population, metric direction, snapshot boundary and unavailable data. Accessible tables mirror every plotted value.


Highest retained Intelligence Index configuration per model; retained snapshot, not current production validation.




Highest retained TAU Banking configuration per model; higher is better; retained snapshot.



Counts unique model cards by provider/benchmark proof type.



Exact public standard-token rates converted with the retained ECB FX. Qwen uses the $\leq 256K$ input tier in this chart; the calculator applies its higher verified tier above 256K.



SECTION 6

Official Public Proofs & Evidence Boundaries

33 usable source snapshots are content-addressed. Model-family identity comes from official providers; benchmark values come from a retained Artificial Analysis API snapshot and are labelled accordingly.

▼ Official provider documentation — current retrieval

Anthropic, OpenAI, Google, DeepSeek, Moonshot/Kimi, xAI, Z.ai, Alibaba Cloud, Cohere and AWS pages were fetched, redirected URLs retained, and SHA-256 hashes recorded. Click any model proof badge in Section 7 to inspect its exact provider link.

► Independent benchmark snapshot — retained, not freshly refreshed

► Pricing and FX boundary

24 Unique Model Cards — 40 Benchmark Configurations

Each model appears once. Adaptive Reasoning and effort labels are mapped inside the card as benchmark configurations, not counted as separate models. Every card explains the score range, evidence and open gaps in simple words. INR values use the ECB 2026-07-24 cross-rate of ₹96.567636/USD.

ANTHROPIC

Provider + benchmark

Claude Opus 5

5 benchmark configurations

One Anthropic model represented by 5 retained benchmark configurations. The configuration labels are grouped here instead of being counted as separate model cards.

0.2330–0.3278

Finance score range

43.9–63.8 tps

Generation speed range

₹482.84 (\$5)

Input price / 1M tokens

₹2414.19 (\$25)

Output price / 1M tokens

Why compare it: Use the ranges to see how benchmark settings changed the result. Open the guide to compare all 5 configurations side by side.

View 5 configurations + evidence

Retained benchmark rate; exact SKU price needs claim-level verification

ANTHROPIC

Provider + benchmark

Claude Fable 5

1 benchmark configuration

An Anthropic model row. The provider documents this model family, and the retained independent benchmark snapshot contains the exact row name. The retained benchmark row uses the setting: Adaptive Reasoning, Max Effort, Opus 4.8 Fallback.

0.2680

Finance score

58.3 tps

Generation speed

₹965.68 (\$10)

Input price / 1M tokens

₹4828.38 (\$50)

Output price / 1M tokens

Why compare it: Finance comparison: its retained TAU Banking score is 0.2680. Treat this as a shortlist signal, not production proof.

[View 1 configuration + evidence](#)

Retained benchmark rate; exact SKU price needs claim-level verification

OPENAI

[Provider + benchmark](#)

GPT-5.6 Sol

5 benchmark configurations

One OpenAI model represented by 5 retained benchmark configurations. The configuration labels are grouped here instead of being counted as separate model cards.

0.2440–0.3299

Finance score range

60.0–73.9 tps

Generation speed range

₹482.84 (\$5)

Input price / 1M tokens

₹2897.03 (\$30)

Output price / 1M tokens

Why compare it: Use the ranges to see how benchmark settings changed the result. Open the guide to compare all 5 configurations side by side.

[View 5 configurations + evidence](#)

Retained benchmark rate; exact SKU price needs claim-level verification

KIMI

[Provider + benchmark](#)

Kimi K3

1 benchmark configuration

A Kimi model row. The provider documents this model family, and the retained independent benchmark snapshot contains the exact row name.

0.3340

Finance score

33.1 tps

Generation speed

₹289.70 (\$3)

Input price / 1M tokens

₹1448.51 (\$15)

Output price / 1M tokens

Why compare it: Finance comparison: its retained TAU Banking score is 0.3340. Treat this as a shortlist signal, not production proof.

[View 1 configuration + evidence](#)

Retained benchmark rate; exact SKU price needs claim-level verification

ANTHROPIC

[Provider + benchmark](#)

Claude Opus 4.8

1 benchmark configuration

An Anthropic model row. The provider documents this model family, and the retained independent benchmark snapshot contains the exact row name. The retained benchmark row uses the setting: Adaptive Reasoning, Max Effort.

0.2763

Finance score

62.5 tps

Generation speed

₹482.84 (\$5)

Input price / 1M tokens

₹2414.19 (\$25)

Output price / 1M tokens

Why compare it: Finance comparison: its retained TAU Banking score is 0.2763. Treat this as a shortlist signal, not production proof.

[View 1 configuration + evidence](#)

Retained benchmark rate; exact SKU price needs claim-level verification

OPENAI

[Provider + benchmark](#)

GPT-5.6 Terra

4 benchmark configurations

One OpenAI model represented by 4 retained benchmark configurations. The configuration labels are grouped here instead of being counted as separate model cards.

0.1938–0.3175

Finance score range

114.6–128.0 tps

Generation speed range

₹241.42 (\$2.5)

Input price / 1M tokens

₹1448.51 (\$15)

Output price / 1M tokens

Why compare it: Use the ranges to see how benchmark settings changed the result. Open the guide to compare all 4 configurations side by side.

[View 4 configurations + evidence](#)

Retained benchmark rate; exact SKU price needs claim-level verification

OPENAI

[Provider + benchmark](#)

GPT-5.5

3 benchmark configurations

One OpenAI model represented by 3 retained benchmark configurations. The configuration labels are grouped here instead of being counted as separate model cards.

0.2577–0.3134

Finance score range

Not available

Generation speed range

₹482.84 (\$5)

Input price / 1M tokens

₹2897.03 (\$30)

Output price / 1M tokens

Why compare it: Use the ranges to see how benchmark settings changed the result. Open the guide to compare all 3 configurations side by side.

[View 3 configurations + evidence](#)

Retained benchmark rate; exact SKU price needs claim-level verification

SPACE XAI

[Provider + benchmark](#)

Grok 4.5

1 benchmark configuration

A SpaceXAI model row. The provider documents this model family, and the retained independent benchmark snapshot contains the exact row name. The retained benchmark row uses the setting: high.

0.3258

Finance score

55.8 tps

Generation speed

₹193.14 (\$2)

Input price / 1M tokens

₹579.41 (\$6)

Output price / 1M tokens

Why compare it: Finance comparison: its retained TAU Banking score is 0.3258. Treat this as a shortlist signal, not production proof.

[View 1 configuration + evidence](#)

Retained benchmark rate; exact SKU price needs claim-level verification

ANTHROPIC

Provider + benchmark

Claude Opus 4.7

1 benchmark configuration

An Anthropic model row. The provider documents this model family, and the retained independent benchmark snapshot contains the exact row name. The retained benchmark row uses the setting: Adaptive Reasoning, Max Effort.

0.2887 Finance score	Not available Generation speed
₹482.84 (\$5) Input price / 1M tokens	₹2414.19 (\$25) Output price / 1M tokens

Why compare it: Finance comparison: its retained TAU Banking score is 0.2887. Treat this as a shortlist signal, not production proof.

View 1 configuration + evidence

Retained benchmark rate; exact SKU price needs claim-level verification

ANTHROPIC

Provider + benchmark

Claude Sonnet 5

1 benchmark configuration

An Anthropic model row. The provider documents this model family, and the retained independent benchmark snapshot contains the exact row name. The retained benchmark row uses the setting: Adaptive Reasoning, Max Effort.

0.2825 Finance score	83.4 tps Generation speed
₹193.14 (\$2) Input price / 1M tokens	₹965.68 (\$10) Output price / 1M tokens

Why compare it: Finance comparison: its retained TAU Banking score is 0.2825. Treat this as a shortlist signal, not production proof.

View 1 configuration + evidence

Retained benchmark rate; exact SKU price needs claim-level verification

OPENAI

Provider + benchmark

GPT-5.4

1 benchmark configuration

An OpenAI model row. The provider documents this model family, and the retained independent benchmark snapshot contains the exact row name. The retained benchmark row uses the setting: xhigh.

0.3031 Finance score	Not available Generation speed
₹241.42 (\$2.5) Input price / 1M tokens	₹1448.51 (\$15) Output price / 1M tokens

Why compare it: Finance comparison: its retained TAU Banking score is 0.3031. Treat this as a shortlist signal, not production proof.

[View 1 configuration + evidence](#)

Retained benchmark rate; exact SKU price needs claim-level verification

OPENAI

[Provider + benchmark](#)

GPT-5.6 Luna

3 benchmark configurations

One OpenAI model represented by 3 retained benchmark configurations. The configuration labels are grouped here instead of being counted as separate model cards.

0.2227–0.2722 Finance score range	161.2–174.3 tps Generation speed range
₹96.57 (\$1) Input price / 1M tokens	₹579.41 (\$6) Output price / 1M tokens

Why compare it: Use the ranges to see how benchmark settings changed the result. Open the guide to compare all 3 configurations side by side.

[View 3 configurations + evidence](#)

Retained benchmark rate; exact SKU price needs claim-level verification

Z AI

[Provider + benchmark](#)

GLM-5.2

1 benchmark configuration

A Z AI model row. The provider documents this model family, and the retained independent benchmark snapshot contains the exact row name. The retained benchmark row uses the setting: max.

0.2680

Finance score

156.7 tps

Generation speed

₹135.19 (\$1.4)

Input price / 1M tokens

₹424.90 (\$4.4)

Output price / 1M tokens

Why compare it: Finance comparison: its retained TAU Banking score is 0.2680. Treat this as a shortlist signal, not production proof.

[View 1 configuration + evidence](#)

Retained benchmark rate; exact SKU price needs claim-level verification

META

Benchmark only

Muse Spark 1.1

1 benchmark configuration

A model name found in the retained independent benchmark snapshot, but current official provider documentation was not established. The retained benchmark row uses the setting: xhigh.

0.2515

Finance score

123.9 tps

Generation speed

₹120.71 (\$1.25)

Input price / 1M tokens

₹410.41 (\$4.25)

Output price / 1M tokens

Why compare it: Finance comparison: its retained TAU Banking score is 0.2515. Treat this as a shortlist signal, not production proof.

[View 1 configuration + evidence](#)

Retained benchmark rate; exact SKU price needs claim-level verification

GOOGLE

Provider + benchmark

Gemini 3.5 Flash

2 benchmark configurations

One Google model represented by 2 retained benchmark configurations. The configuration labels are grouped here instead of being counted as separate model cards.

0.2536

Finance score range

250.3–265.2 tps

Generation speed range

₹144.85 (\$1.5)

Input price / 1M tokens

₹869.11 (\$9)

Output price / 1M tokens

Why compare it: Use the ranges to see how benchmark settings changed the result. Open the guide to compare all 2 configurations side by side.

[View 2 configurations + evidence](#)

Retained benchmark rate; exact SKU price needs claim-level verification

GOOGLE

[Provider + benchmark](#)

Gemini 3.6 Flash

1 benchmark configuration

A Google model row. The provider documents this model family, and the retained independent benchmark snapshot contains the exact row name. The retained benchmark row uses the setting: high.

0.2454

Finance score

243.9 tps

Generation speed

₹144.85 (\$1.5)

Input price / 1M tokens

₹724.26 (\$7.5)

Output price / 1M tokens

Why compare it: Speed comparison: the retained snapshot reports 243.9 output tokens per second. Useful when throughput matters.

[View 1 configuration + evidence](#)

Retained benchmark rate; exact SKU price needs claim-level verification

COHERE

[Provider only](#)

Cohere Command A

1 benchmark configuration

A Cohere model documented by the provider. This report does not have an exact matching row in the retained independent benchmark snapshot.

0.2850

Finance score

112.0 tps

Generation speed

₹48.28 (\$0.5)

Input price / 1M tokens

₹144.85 (\$1.5)

Output price / 1M tokens

Why compare it: Finance comparison: its retained TAU Banking score is 0.2850. Treat this as a shortlist signal, not production proof.

[View 1 configuration + evidence](#)

Retained benchmark rate; exact SKU price needs claim-level verification

ANTHROPIC

[Provider + benchmark](#)

Claude Sonnet 4.6

1 benchmark configuration

An Anthropic model row. The provider documents this model family, and the retained independent benchmark snapshot contains the exact row name. The retained benchmark row uses the setting: Adaptive Reasoning, Max Effort.

0.3052

Finance score

Not available

Generation speed

₹289.70 (\$3)

Input price / 1M tokens

₹1448.51 (\$15)

Output price / 1M tokens

Why compare it: Finance comparison: its retained TAU Banking score is 0.3052. Treat this as a shortlist signal, not production proof.

[View 1 configuration + evidence](#)

Retained benchmark rate; exact SKU price needs claim-level verification

AMAZON

[Provider only](#)

Amazon Nova Pro

1 benchmark configuration

An Amazon model documented by the provider. This report does not have an exact matching row in the retained independent benchmark snapshot.

0.2710

Finance score

145.0 tps

Generation speed

₹77.25 (\$0.8)

Input price / 1M tokens

₹309.02 (\$3.2)

Output price / 1M tokens

Why compare it: Finance comparison: its retained TAU Banking score is 0.2710. Treat this as a shortlist signal, not production proof.

[View 1 configuration + evidence](#)

GOOGLE

Provider + benchmark

Gemini 3.1 Pro Preview

1 benchmark configuration

A Google model row. The provider documents this model family, and the retained independent benchmark snapshot contains the exact row name.

0.1649 Finance score	132.2 tps Generation speed
₹193.14 (\$2) Input price / 1M tokens	₹1158.81 (\$12) Output price / 1M tokens

Why compare it: Use it as a cost and benchmark comparison row. No production recommendation is implied.

View 1 configuration + evidence

ALIBABA

Provider + benchmark

Qwen3.7 Max

1 benchmark configuration

An Alibaba model row. The provider documents this model family, and the retained independent benchmark snapshot contains the exact row name.

0.1093 Finance score	199.6 tps Generation speed
₹241.42 (\$2.5) Input price / 1M tokens	₹724.26 (\$7.5) Output price / 1M tokens

Why compare it: Speed comparison: the retained snapshot reports 199.6 output tokens per second. Useful when throughput matters.

View 1 configuration + evidence

KIMI

Provider + benchmark

Kimi K2.6

1 benchmark configuration

A general-purpose Kimi model. The provider documents text, image and video input, thinking modes, tool calls and a 262,144-token context window.

0.2062 Finance score	Not available Generation speed
₹91.74 (\$0.95) Input price / 1M tokens	₹386.27 (\$4) Output price / 1M tokens

Why compare it: Useful as a lower-cost finance and agent-workflow candidate. Its retained TAU Banking score is 0.2062.

[View 1 configuration + evidence](#)

Verified exact official rate card

ALIBABA

[Provider + benchmark](#)

Qwen3.7 Plus

1 benchmark configuration

A lower-cost Qwen3.7 API tier listed by Alibaba Cloud Model Studio with the exact model ID qwen3.7-plus.

0.1787 Finance score	52.8 tps Generation speed
₹38.63 (\$0.4) Input price / 1M tokens	₹154.51 (\$1.6) Output price / 1M tokens

Why compare it: Useful when price matters more than using the Max tier. Its retained TAU Banking score is 0.1787.

[View 1 configuration + evidence](#)

Verified exact official rate card

GOOGLE

[Provider + benchmark](#)

Gemini 3.5 Flash-Lite

1 benchmark configuration

Google describes Gemini 3.5 Flash-Lite as its fastest and most cost-effective 3.5 model for high-throughput execution.

0.1649

Finance score

362.2 tps

Generation speed

₹28.97 (\$0.3)

Input price / 1M tokens

₹241.42 (\$2.5)

Output price / 1M tokens

Why compare it: Useful for high-volume work. The retained benchmark snapshot reports 362.2 output tokens per second.

[View 1 configuration + evidence](#)

Verified exact official rate card

Enterprise Governance, SEBI & Global Compliance

SEBI Investment Adviser Circulars, FINRA Rule 3110 / 2010, SEC Rule 204-2, EU AI Act Annex III / Article 15, SOC 2 Type II, and FedRAMP High define distinct governance boundaries. This section separates legal scope, attestations, audit logs, and SARVAX engineering precision policy.

SEBI IA Regulations

SEBI Investment Adviser & Advisory Compliance

Under SEBI (Investment Advisers) Regulations, 2013 and SEBI Cybersecurity Framework, AI systems generating portfolio advice or mutual fund recommendations must maintain deterministic suitability checks, risk-profiling validation, and mandatory advisor sign-off before client dispatch. Models cannot directly execute transactions without human authorization.

[SEBI IA Regulations](#)

FINRA & SEC Audit Trails

FINRA Rule 3110 / 2010 & SEC Rule 204-2

FINRA Rule 3110 (Supervision) and SEC Rule 204-2 (Books and Records) require complete immutability and 7-year retention for all AI-generated client communications, meeting notes, and suitability logs. SARVAX enforces structured JSON schema audit logging with hash signatures for every agent call.

[FINRA Rule 3110](#)

► SARVAX Indian Wealth Advisory Compliance Pack

Verified legal text

Annex III applicability is use-case specific

Annex III includes AI used to evaluate the creditworthiness or credit score of natural persons and AI used for risk assessment and pricing for natural persons in life and health insurance. A portfolio-analysis or wealth-advisory system is not automatically Annex III high-risk solely because it operates in finance; intended purpose and material influence must be assessed.

[Official EU Annex III](#)

Verified legal text

Article 15: accuracy, robustness and cybersecurity

Article 15 requires high-risk AI systems to achieve an appropriate level of accuracy, robustness and cybersecurity throughout their lifecycle, declare relevant accuracy metrics and be resilient to errors, faults and attacks. It does not name INT4, FP8 or BF16.

[European Commission Article 15](#)

► SARVAX evidence required before EU deployment

Third-party attestation

SOC 2 Type II

SOC 2 Type II concerns controls for a defined service-organization system over a review period. It is not a compliance badge inherited by an individual model. Review the report scope, Trust Services Criteria, examination period, exceptions and bridge coverage.

[AICPA SOC resources](#)

Contractual mechanism

HIPAA Business Associate Agreement

Where PHI is handled, a covered entity must obtain appropriate assurances through a business-associate contract and confirm the relevant service and data flow are in scope. "HIPAA eligible" model availability alone is insufficient.

[HHS business-associate guidance](#)

► Evidence required

AWS statement verified

Amazon Bedrock in GovCloud

AWS states that Amazon Bedrock is a FedRAMP High authorized service in AWS GovCloud (US-West). That statement applies to the authorized cloud service and boundary—not generically to every model, application architecture or customer deployment.

[AWS Bedrock security & compliance](#)

100% sovereignty claim rejected

Private VPC is not automatic air-gapping

Self-hosting a model in a VPC does not by itself guarantee air-gapping or “100% cloud data sovereignty.” Sovereignty depends on region, identity, egress, private endpoints, encryption keys, logs, backups, telemetry, support access, subprocessors and the model-artifact supply chain.

[AWS GovCloud](#)

► FedRAMP High review boundary

Not statutory text

Article 15 does not prohibit INT4

The statement “INT4 is strictly prohibited and FP8/BF16 are required” is not present in Article 15. Quantization can affect numeric fidelity, but the legal requirement is outcome-based accuracy, robustness and cybersecurity.

[Verify Article 15](#)

SARVAX engineering policy

Material financial-risk scoring precision gate

SARVAX policy: do not deploy INT4 for material financial-risk scoring unless model-specific validation demonstrates no material degradation against the approved reference precision. FP8/BF16 may be approved only after the same tests; neither is automatically safe.

► **Required precision-validation pack**

SARVAX AI Model Intelligence & Enterprise Architecture Report — July 2026 Edition

Provider documentation retrieved 2026-07-25T12:56:15+00:00 · ECB FX reference 2026-07-24 · retained Artificial
Analysis benchmark snapshot
C3A Labs R&D