Evidence-scoped catalog of 24 model cards across 40 retained benchmark configurations. Proposed routing is evaluated separately and may include provider-verified models without catalog-grade benchmark matching.

Evidence:

FX BASIS

PROOF STATE

Verdict: Conditional Hybrid Cascade — Evidence Gates Remain

A hybrid cascade remains a valid **architecture pattern**, not a verified model-selection decision. Twenty-one unique models have both provider-family documentation and at least one exact retained benchmark row; three need more proof. Financial savings, provider rate limits and production SLAs require exact SKU, contract and trace evidence before GTM approval.

Decision: architecture exploration may continue; procurement, financial commitment and regulated production routing remain conditional.

Evidence boundary: provider pages establish model-family documentation. Artificial Analysis metrics are from a retained API snapshot that could not be freshly refreshed because the API credential was unavailable. A benchmark row does not establish provider availability, SLA, compliance or customer access.

21	**3**	**12**	**3**
DUAL-PROOF MODEL CARDS	WEALTH-ADVISORY JOURNEYS	DOMAIN AGENT ROLES	PENDING VERIFICATION

SECTION 1

Executive Product Verdict: Wealth Advisory Architecture

A **Product and Enterprise Margins POV** evaluating the frontier landscape for long-running financial wealth workflows.

SARVAX WEALTH-ADVISORY ARCHITECTURE

Models underneath the product—not the product itself

12 agents · 3 advisory journeys

Wealth advisory work arrives through meetings, CRM history, email, WhatsApp, holdings, market data, KYC records, tasks and client documents. SARVAX owns the journey: intake, agent orchestration, memory, approvals, integrations and writeback. Foundation models perform bounded routing, context assembly, vision and reasoning roles underneath it.

> **Product recommendation:** Design around advisor journeys—not PDF types. Keep deterministic calculations, policy enforcement, authorization and human approval outside the model chain.

Relationship intake

Meeting Assistant + CRM/MCP context

↓

The Router

Gemini 3.5 Flash-Lite

↓

The Reader

DeepSeek V4 Pro

↓

Relationship agents

Pre/Post Meeting + Client Strategist

Advisor approval

Workflow 2.0 approval gate

↓

System writeback

CRM, tasks, notifications, artifacts

Turns meetings and ongoing client communication into advisor-ready context, follow-ups and controlled system updates.

SELECTED SARVAX JOURNEY

Client Relationship Intelligence

ADVISOR OUTCOME

Every advisor starts prepared, leaves with structured actions and keeps the client record current.

SARVAX AGENTS

- Pre-Meeting Analysis Agent
- Post-Meeting Analysis Agent
- Client Intelligence Strategist

PRODUCT SURFACES

- Meeting Assistant
- Agent Builder
- Agent Memory
- Workflow 2.0
- MCP / integrations layer
- Notifications and writeback

INPUTS ACROSS THE JOURNEY

- Meeting Assistant transcripts
- Email and WhatsApp context
- CRM household history
- Advisor notes and tasks

CONTROL BOUNDARY

Meeting intelligence drafts context and actions. The advisor approves material client communication and every financial recommendation.

BOUNDED MODEL ROLES UNDERNEATH THESE JOURNEYS

CORE STAGE 2

Moonshot AI

Product role: Produces bounded portfolio, tax and compliance analysis drafts inside SARVAX agents. Deterministic calculators and advisor approval own the final decision.

Why: Its retained TAU Banking result is 0.3340, the highest value in this report's retained 40-configuration comparison set. This is a shortlist signal, not proof of production accuracy.

Boundary: Use only where a wealth-advisory journey needs deeper financial reasoning. Meeting notes, routing and routine extraction should not pay this cost or inherit this latency.

0.3340
TAU Banking
Retained independent snapshot

33.1 tps
Speed
Retained independent snapshot

₹289.70
Input / 1M
Retained benchmark rate; exact SKU price needs claim-level verification

---

Task testing required

Provider proof • Benchmark

CORE STAGE 1

DeepSeek

Product role: Normalizes high-volume text from meeting transcripts, email, WhatsApp exports, CRM notes, research and text-native statements before SARVAX agents reason over it.

Why: The official rate card verifies the exact SKU, low cache-hit pricing, JSON output, tool calls and a 1M context window—useful for large client-context assembly.

Boundary: It is a text-context worker, not the product. Image-origin evidence still needs the Vision Gate; extracted facts must retain source references and pass journey-specific validation.

0.2371–0.2577
TAU Banking range
Retained independent configurations

31.2–44.3
General score range
Retained independent configurations

₹42.01
Cache-miss input / 1M
Official DeepSeek rate card

Official SKU and price verified; task quality unproven

Provider proof • Benchmark

CONDITIONAL GATE

Google

Product role: Identifies the wealth-advisory journey, urgency, household context and required SARVAX agent/workflow before expensive processing begins.

Why: Google describes it as its fastest, most cost-effective 3.5 model for high-throughput execution. The retained snapshot reports 362.2 tps; its verified input rate is ₹28.97/1M.

Boundary: Use for routing, classification and low-risk extraction checks—not investment, tax or compliance conclusions. Low-confidence routing must fall back to review.

362.2 tps Speed Retained independent snapshot	**₹28.97** Input / 1M Verified official rate card	**0.1649** TAU Banking Retained independent snapshot

Conditional gate; task testing required

Provider proof • Benchmark

CONDITIONAL GATE

Google

Product role: Handles image-origin evidence—scanned KYC records, photographed forms, charts and statement pages—when a broader wealth-advisory journey actually needs it.

Why: Google's model page describes Gemini 3.6 Flash as supporting agentic and multimodal tasks. The retained snapshot reports 243.9 tps.

Boundary: Multimodal identity does not prove financial-document OCR accuracy. Validate tables, decimal placement, signatures and low-quality scans; escalate low-confidence evidence.

243.9 tps Speed Retained independent snapshot	**50.1** General score Retained independent snapshot	**Exact SKU rate not established** Price Retained benchmark rate; exact SKU price needs claim-level verification

Multimodal identity verified; OCR accuracy unproven

Provider proof • Benchmark

▶ Evidence boundaries and QA corrections

SECTION 2

SARVAX Wealth-Advisory Agent Operating Matrix

Twelve product agents mapped across three advisor journeys. No invented time-saved, accuracy, SLA or unit-economics claims are shown. Economics stay in the scenario calculator until measured traces exist.

CLIENT RELATIONSHIP INTELLIGENCE

Turns meetings and ongoing client communication into advisor-ready context, follow-ups and controlled system updates.

Pre-Meeting Analysis Agent · Post-Meeting Analysis Agent · Client Intelligence Strategist

PRODUCT SURFACES	Meeting Assistant · Agent Builder · Agent Memory · Workflow 2.0 · MCP / integrations layer · Notifications and writeback
INPUTS	Meeting Assistant transcripts · Email and WhatsApp context · CRM household history · Advisor notes and tasks
CONTROL BOUNDARY	Meeting intelligence drafts context and actions. The advisor approves material client communication and every financial recommendation.

EVIDENCE STATUS

Frontend product contracts observed. Backend execution, connector availability, production reliability and branch freshness remain unproven.

PORTFOLIO & MARKET INTELLIGENCE

Combines holdings, market context and client objectives into review-ready portfolio and tax-analysis drafts.

Portfolio Analyst · Market Intelligence Agent

PRODUCT SURFACES	MCP / integrations layer · OneChat projects and artifacts · Agent Builder · Skills · Workflow 2.0 · Approval and audit trail
INPUTS	Holdings and transactions · Market and research feeds · Risk profile and goals · Statements and tax-lot data
CONTROL BOUNDARY	Models can draft analysis. Deterministic portfolio/tax engines validate the math; the advisor owns suitability and final action.

EVIDENCE STATUS

Frontend product contracts observed. Backend execution, connector availability, production reliability and branch freshness remain unproven.

COMPLIANCE & ADVISOR OPERATIONS

Routes KYC, review cadence, pipeline, workload and document events through bounded agents and approval-controlled workflows.

Cadence Ops Lead · Campaign Audience Architect · Pipeline & Escalation Tracker · Team Performance & Workload Insights Agent · KYC Intelligence Officer · KYC/CDD QA Auditor · Document Intelligence & Auto-Population Agent

PRODUCT SURFACES	Agent Builder · Agent Memory · Workflow 2.0 parallel groups · Approval gates · Skills · MCP / integrations layer · Notifications
INPUTS	KYC/CDD records · Client and account changes · Review cadence and pipeline events · Forms, messages and team tasks
CONTROL BOUNDARY	Policy rules and deterministic validation remain outside the model chain. Compliance exceptions and material account changes require named human approval.

EVIDENCE STATUS

Frontend product contracts observed. Backend execution, connector availability, production reliability and branch freshness remain unproven.

SECTION 3

Standard-Token Cost Scenario

Compares standard token rates across frontier models in the repository catalog using exact Decimal arithmetic and the official ECB USD/INR reference rate. Workload volumes are editable planning assumptions—not production invoices or observed SARVAX usage.

Verified inputs

Catalog rate cards and ECB 2026-07-24 USD/INR cross-rate.

Scenario inputs

Input tokens, output tokens and monthly execution volume.

Excluded

Cache, batch, retries, tools, storage, platform fees, taxes and negotiated contracts.

Formula: ((input tokens × input USD/1M) + (output tokens × output USD/1M)) × USD/INR. Model-specific context and pricing tiers are enforced before a result is shown.

Planning inputs

SCENARIO PRESET

Portfolio Review (75,000 in / 8,000 out)

MODEL A

Gemini 3.5 Flash-Lite (Google) — $0.30 in / $2.50 out per 1M [Verified]

MODEL B

DeepSeek V4 Pro (DeepSeek) — $0.435 in / $0.87 out per 1M [Verified]

Input tokens per run **75,000**

Output tokens per run **8,000**

Monthly execution volume **10,000 runs**

FX basis: ECB 2026-07-24 — ₹96.567636/USD. Latest business-day official cross-rate.

Calculated with ECB 2026-07-24 FX and verified public rate inputs. Gemini 3.5 Flash-Lite: Global standard inference; official Google rate card. DeepSeek V4 Pro: Standard cache-miss input and output; official DeepSeek rate card.

Calculated standard-token cost

Global standard inference; official Google rate card

PER RUN

₹4.10

MONTHLY

₹0.4104 lakh / month

PER 1,000 RUNS

₹4104.12

Standard cache-miss input and output; official DeepSeek rate card

PER RUN

₹3.82

MONTHLY

₹0.3823 lakh / month

PER 1,000 RUNS

₹3822.63

CALCULATED PRICE DIFFERENCE

DeepSeek V4 Pro is lower for these planning inputs. This is a rate-only difference—not measured savings.

ANNUALIZED DIFFERENCE

INTERPRETATION

SECTION 4

Wealth-Advisory Control-Flow Sandbox

Tests where product orchestration, deterministic gates and advisor approval must sit before a POC is approved. It demonstrates control order—not production execution quality.

SELECT WEALTH-ADVISORY JOURNEY

Client Relationship Intelligence

Scenario progress — Ready

Evidence boundary Stage sequence only · **No measured tokens, cost, latency, cache, retries or SLA**

DECISION UNDER TEST

Can SARVAX preserve evidence, ownership and human decision rights across the complete meeting journey?

From permitted client context to advisor-approved follow-up and controlled system writeback.

PRODUCT OWNER

Meeting Assistant + Workflow 2.0

DETERMINISTIC OWNER

Policy, authorization and write-integrity services

HUMAN DECISION RIGHT

Named advisor

EXIT GATE

Release only after consent, tenant isolation, source lineage, approval and writeback controls pass representative trace tests.

RELEASE EVIDENCE STILL REQUIRED

Connector execution, persistence, retry behavior and production reliability are not yet proven with representative backend traces.

Control-flow chain Governance and approval scenario — not production telemetry

Pending · Active · Passed

1 **Permitted client-context intake** PRODUCT — PENDING

Calendar, CRM and approved communication context are collected before preparation.

Required gate: Required: consent, tenant access, source timestamps and retention policy.

2 **Pre-Meeting Analysis Agent** AGENT — PENDING

Drafts a source-linked preparation brief and explicit data gaps.

Required gate: Required: unsupported client facts blocked; every material claim links to retained evidence.

3 **Meeting capture and artifact finalization** PRODUCT — PENDING

Meeting Assistant creates the permitted transcript and meeting artifacts.

Required gate: Required: participant consent, speaker attribution review and artifact access policy.

4 **Post-Meeting Analysis Agent** AGENT PENDING

Drafts facts, decisions, risks and action items from the retained artifact.

Required gate: Required: transcript facts remain separate from model interpretation.

5 **Client Intelligence Strategist** AGENT PENDING

Drafts next-best-action options for advisor review.

Required gate: Required: no autonomous product recommendation, promise or client communication.

6 **Advisor decision gate** HUMAN PENDING

The named advisor accepts, edits or rejects facts and proposed actions.

Required gate: Required before any client communication or material CRM/task update.

7 **Controlled Workflow 2.0 writeback** TOOL PENDING

Only approved records are submitted to authorized connected systems.

Required gate: Release evidence required: idempotency key, immutable audit event, retry/dead-letter policy and compensating action.

SECTION 5

Four Evidence-Backed Comparison Charts

Charts read the central runtime dataset. Each chart discloses population, metric direction, snapshot boundary and unavailable data. Accessible tables mirror every plotted value.

Highest retained Intelligence Index configuration per model; retained snapshot, not current production validation.

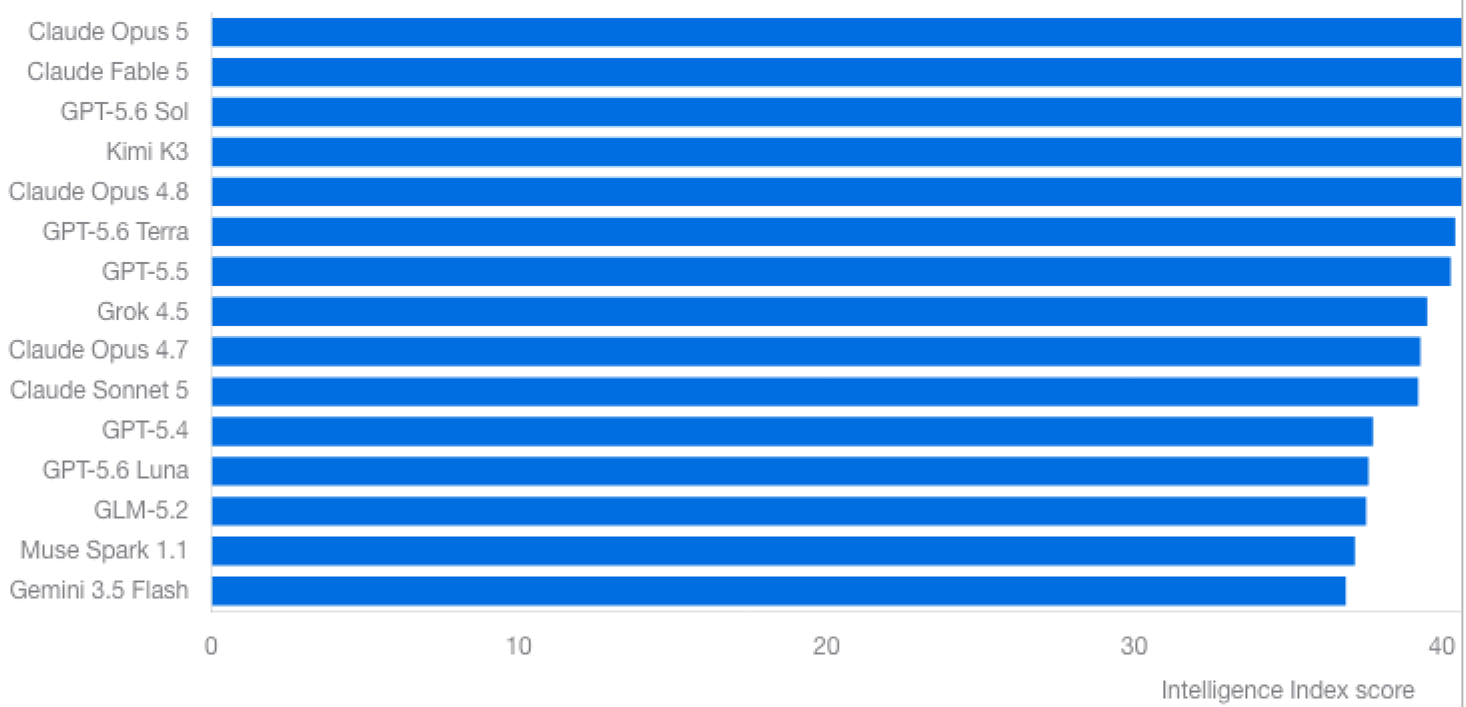

Highest retained TAU Banking configuration per model; higher is better; retained snapshot.

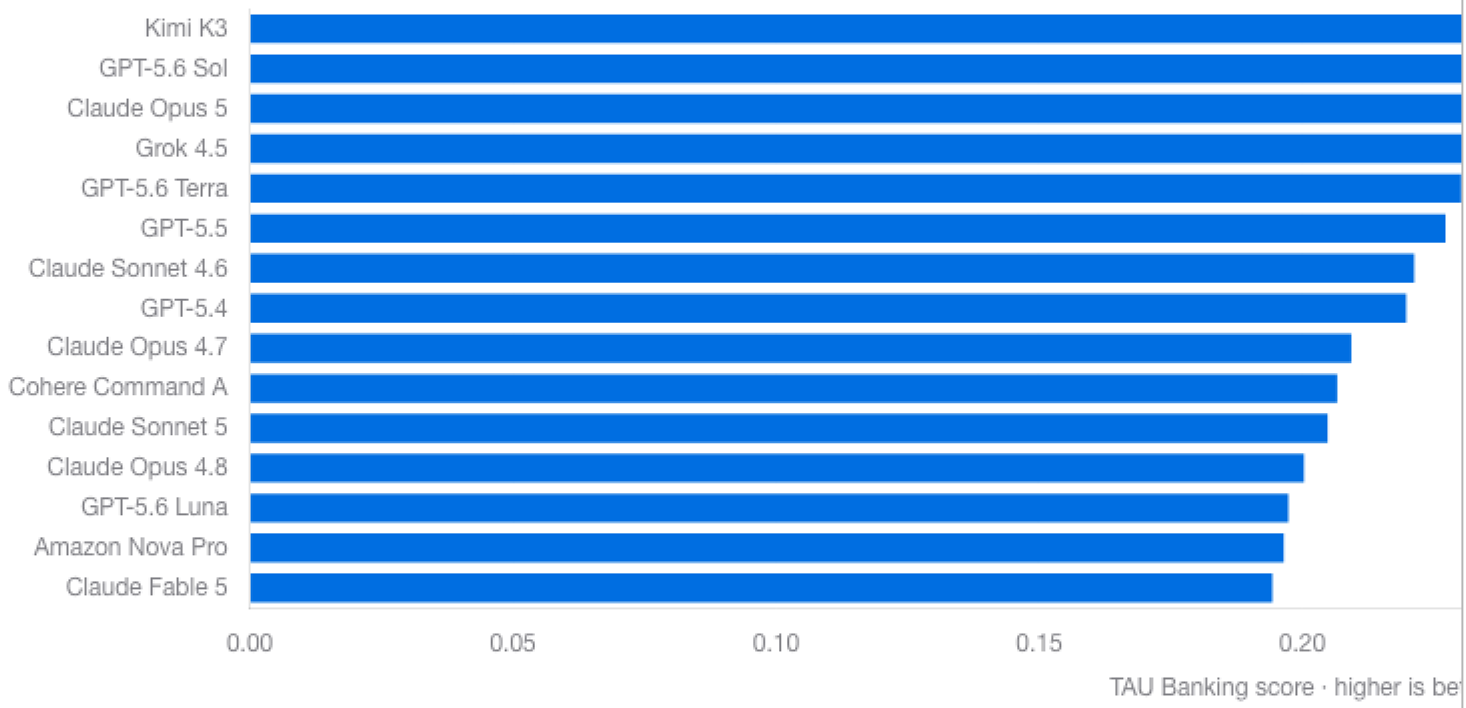

Counts unique model cards by provider/benchmark proof type.

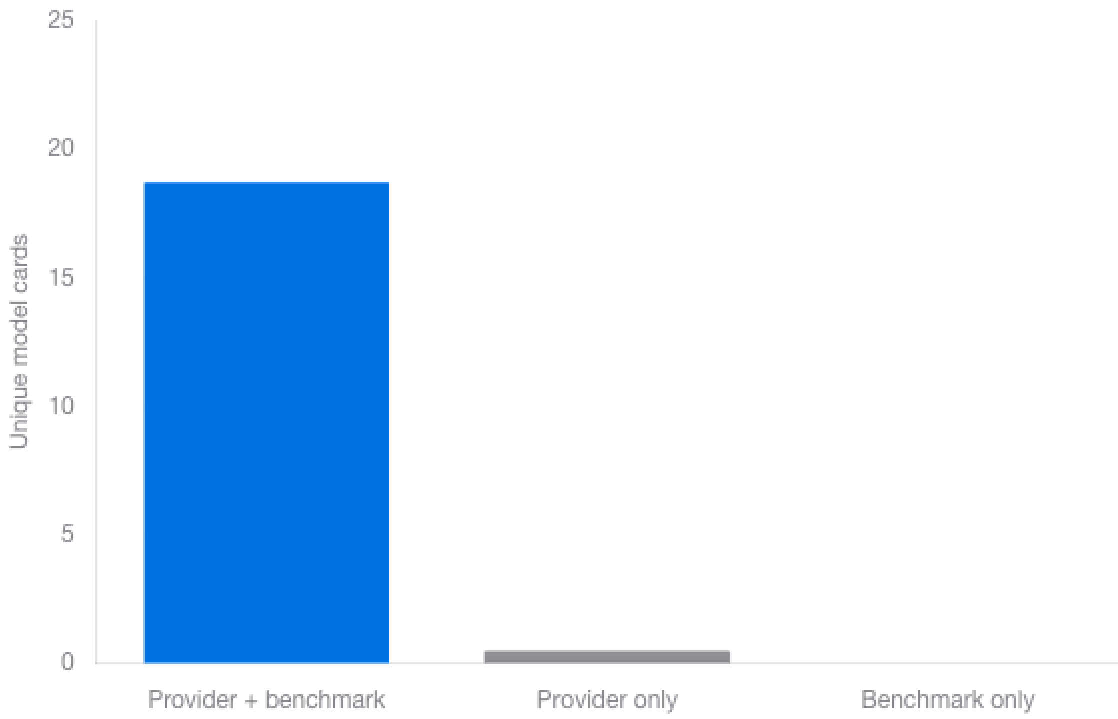

Exact public standard-token rates converted with the retained ECB FX. Qwen uses the ≤256K input tier in this chart; the calculator applies its higher verified tier above 256K.

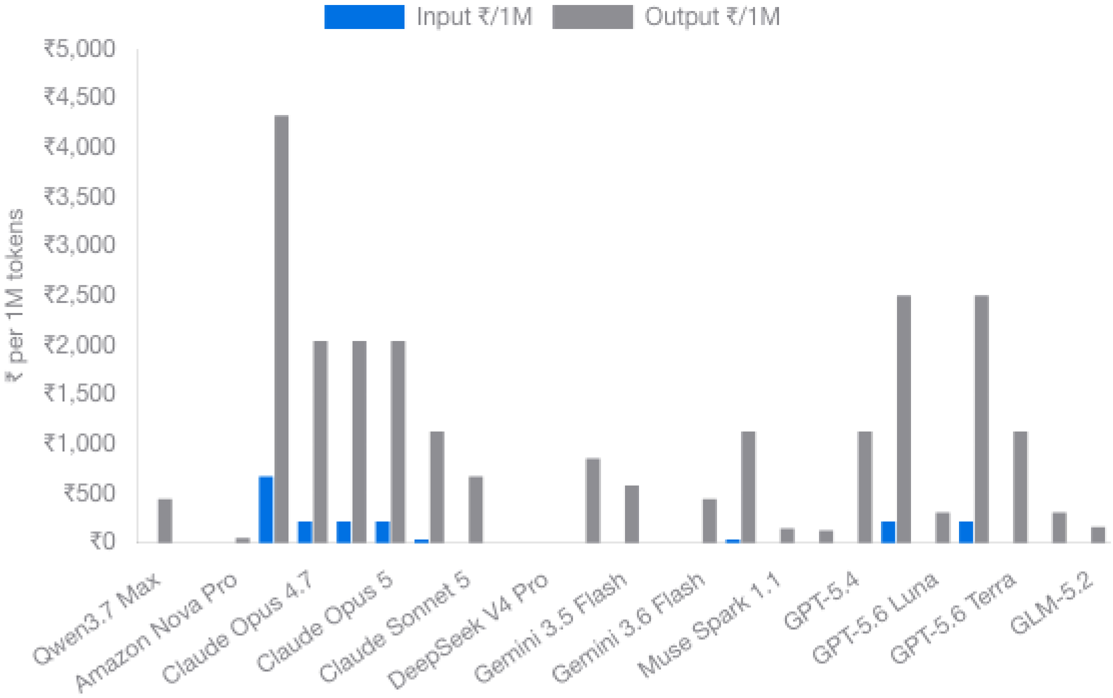

SECTION 6

Official Public Proofs & Evidence Boundaries

33 usable source snapshots are content-addressed. Model-family identity comes from official providers; benchmark values come from a retained Artificial Analysis API snapshot and are labelled accordingly.

▼ **Official provider documentation — current retrieval**

Anthropic, OpenAI, Google, DeepSeek, Moonshot/Kimi, xAI, Z.ai, Alibaba Cloud, Cohere and AWS pages were fetched, redirected URLs retained, and SHA-256 hashes recorded. Click any model proof badge in Section 7 to inspect its exact provider link.

▶ **Independent benchmark snapshot — retained, not freshly refreshed**

▶ **Pricing and FX boundary**

SECTION 7

24 Unique Model Cards — 40 Benchmark Configurations

Each model appears once. Adaptive Reasoning and effort labels are mapped inside the card as benchmark configurations, not counted as separate models. Every card explains the score range, evidence and open gaps in simple words. INR values use the ECB 2026-07-24 cross-rate of ₹96.567636/USD.

ANTHROPIC

Claude Opus 5

Provider + benchmark

5 benchmark configurations

One Anthropic model represented by 5 retained benchmark configurations. The configuration labels are grouped here instead of being counted as separate model cards.

0.2330–0.3278
Finance score range · higher is better

50.6–60.7
General score range · higher is better

43.9–63.8 tps
Generation speed range

₹482.84
Input price / 1M tokens

Why compare it: Use the ranges to see how benchmark settings changed the result. Open the guide to compare all 5 configurations side by side.

View 5 configurations + evidence

Retained benchmark rate; exact SKU price needs claim-level verification

ANTHROPIC

Claude Fable 5

Provider + benchmark

1 benchmark configuration

An Anthropic model row. The provider documents this model family, and the retained independent benchmark snapshot contains the exact row name. The retained benchmark row uses the setting: Adaptive Reasoning, Max Effort, Opus 4.8 Fallback.

0.2680
Finance score · higher is better

59.9
General score · higher is better

58.3 tps
Generation speed

₹965.68
Input price / 1M tokens

> **Why compare it:** Finance comparison: its retained TAU Banking score is 0.2680. Treat this as a shortlist signal, not production proof.

View 1 configuration + evidence

Retained benchmark rate; exact SKU price needs claim-level verification

OPENAI

GPT-5.6 Sol

Provider + benchmark

5 benchmark configurations

One OpenAI model represented by 5 retained benchmark configurations. The configuration labels are grouped here instead of being counted as separate model cards.

0.2440–0.3299
Finance score range · higher is better

49.4–58.9
General score range · higher is better

60.0–73.9 tps
Generation speed range

₹482.84
Input price / 1M tokens

> **Why compare it:** Use the ranges to see how benchmark settings changed the result. Open the guide to compare all 5 configurations side by side.

View 5 configurations + evidence

Retained benchmark rate; exact SKU price needs claim-level verification

KIMI

Kimi K3

Provider + benchmark

1 benchmark configuration

A Kimi model row. The provider documents this model family, and the retained independent benchmark snapshot contains the exact row name.

0.3340
Finance score · higher is better

57.1
General score · higher is better

33.1 tps
Generation speed

₹289.70
Input price / 1M tokens

Why compare it: Finance comparison: its retained TAU Banking score is 0.3340. Treat this as a shortlist signal, not production proof.

View 1 configuration + evidence

Retained benchmark rate; exact SKU price needs claim-level verification

ANTHROPIC

Provider + benchmark

Claude Opus 4.8

1 benchmark configuration

An Anthropic model row. The provider documents this model family, and the retained independent benchmark snapshot contains the exact row name. The retained benchmark row uses the setting: Adaptive Reasoning, Max Effort.

0.2763
Finance score · higher is better

55.7
General score · higher is better

62.5 tps
Generation speed

₹482.84
Input price / 1M tokens

Why compare it: Finance comparison: its retained TAU Banking score is 0.2763. Treat this as a shortlist signal, not production proof.

View 1 configuration + evidence

Retained benchmark rate; exact SKU price needs claim-level verification

OPENAI

Provider + benchmark

GPT-5.6 Terra

4 benchmark configurations

One OpenAI model represented by 4 retained benchmark configurations. The configuration labels are grouped here instead of being counted as separate model cards.

0.1938–0.3175
Finance score range · higher is better

45.6–55.0
General score range · higher is better

114.6–128.0 tps
Generation speed range

₹241.42
Input price / 1M tokens

Why compare it: Use the ranges to see how benchmark settings changed the result. Open the guide to compare all 4 configurations side by side.

View 4 configurations + evidence

Retained benchmark rate; exact SKU price needs claim-level verification

OPENAI

GPT-5.5

Provider + benchmark

3 benchmark configurations

One OpenAI model represented by 3 retained benchmark configurations. The configuration labels are grouped here instead of being counted as separate model cards.

0.2577–0.3134
Finance score range · higher is better

50.4–54.8
General score range · higher is better

Not available
Generation speed range

₹482.84
Input price / 1M tokens

Why compare it: Use the ranges to see how benchmark settings changed the result. Open the guide to compare all 3 configurations side by side.

View 3 configurations + evidence

Retained benchmark rate; exact SKU price needs claim-level verification

SPACEXAI

Grok 4.5

Provider + benchmark

1 benchmark configuration

A SpaceXAI model row. The provider documents this model family, and the retained independent benchmark snapshot contains the exact row name. The retained benchmark row uses the setting: high.

0.3258
Finance score · higher is better

53.8
General score · higher is better

55.8 tps
Generation speed

₹193.14
Input price / 1M tokens

Why compare it: Finance comparison: its retained TAU Banking score is 0.3258. Treat this as a shortlist signal, not production proof.

View 1 configuration + evidence

Retained benchmark rate; exact SKU price needs claim-level verification

ANTHROPIC

Claude Opus 4.7

Provider + benchmark

1 benchmark configuration

An Anthropic model row. The provider documents this model family, and the retained independent benchmark snapshot contains the exact row name. The retained benchmark row uses the setting: Adaptive Reasoning, Max Effort.

0.2887
Finance score · higher is better

53.5
General score · higher is better

Not available
Generation speed

₹482.84
Input price / 1M tokens

> **Why compare it:** Finance comparison: its retained TAU Banking score is 0.2887. Treat this as a shortlist signal, not production proof.

View 1 configuration + evidence

Retained benchmark rate; exact SKU price needs claim-level verification

ANTHROPIC

Claude Sonnet 5

Provider + benchmark

1 benchmark configuration

An Anthropic model row. The provider documents this model family, and the retained independent benchmark snapshot contains the exact row name. The retained benchmark row uses the setting: Adaptive Reasoning, Max Effort.

0.2825
Finance score · higher is better

53.4
General score · higher is better

83.4 tps
Generation speed

₹193.14
Input price / 1M tokens

> **Why compare it:** Finance comparison: its retained TAU Banking score is 0.2825. Treat this as a shortlist signal, not production proof.

View 1 configuration + evidence

Retained benchmark rate; exact SKU price needs claim-level verification

OPENAI

GPT-5.4

Provider + benchmark

1 benchmark configuration

An OpenAI model row. The provider documents this model family, and the retained independent benchmark snapshot contains the exact row name. The retained benchmark row uses the setting: xhigh.

0.3031
Finance score · higher is better

51.4
General score · higher is better

Not available
Generation speed

₹241.42
Input price / 1M tokens

Why compare it: Finance comparison: its retained TAU Banking score is 0.3031. Treat this as a shortlist signal, not production proof.

View 1 configuration + evidence

Retained benchmark rate; exact SKU price needs claim-level verification

OPENAI

GPT-5.6 Luna

Provider + benchmark

3 benchmark configurations

One OpenAI model represented by 3 retained benchmark configurations. The configuration labels are grouped here instead of being counted as separate model cards.

0.2227–0.2722
Finance score range · higher is better

46.1–51.2
General score range · higher is better

161.2–174.3 tps
Generation speed range

₹96.57
Input price / 1M tokens

Why compare it: Use the ranges to see how benchmark settings changed the result. Open the guide to compare all 3 configurations side by side.

View 3 configurations + evidence

Retained benchmark rate; exact SKU price needs claim-level verification

Z AI

GLM-5.2

Provider + benchmark

1 benchmark configuration

A Z AI model row. The provider documents this model family, and the retained independent benchmark snapshot contains the exact row name. The retained benchmark row uses the setting: max.

0.2680
Finance score · higher is better

51.1
General score · higher is better

156.7 tps
Generation speed

₹135.19
Input price / 1M tokens

Why compare it: Finance comparison: its retained TAU Banking score is 0.2680. Treat this as a shortlist signal, not production proof.

View 1 configuration + evidence

Retained benchmark rate; exact SKU price needs claim-level verification

META

Benchmark only

Muse Spark 1.1

1 benchmark configuration

A model name found in the retained independent benchmark snapshot, but current official provider documentation was not established. The retained benchmark row uses the setting: xhigh.

0.2515
Finance score · higher is better

50.6
General score · higher is better

123.9 tps
Generation speed

₹120.71
Input price / 1M tokens

Why compare it: Finance comparison: its retained TAU Banking score is 0.2515. Treat this as a shortlist signal, not production proof.

View 1 configuration + evidence

Retained benchmark rate; exact SKU price needs claim-level verification

GOOGLE

Provider + benchmark

Gemini 3.5 Flash

2 benchmark configurations

One Google model represented by 2 retained benchmark configurations. The configuration labels are grouped here instead of being counted as separate model cards.

0.2536
Finance score range · higher is better

45.4–50.2
General score range · higher is better

250.3–265.2 tps
Generation speed range

₹144.85
Input price / 1M tokens

Why compare it: Use the ranges to see how benchmark settings changed the result. Open the guide to compare all 2 configurations side by side.

View 2 configurations + evidence

Retained benchmark rate; exact SKU price needs claim-level verification

GOOGLE

Gemini 3.6 Flash

Provider + benchmark

1 benchmark configuration

A Google model row. The provider documents this model family, and the retained independent benchmark snapshot contains the exact row name. The retained benchmark row uses the setting: high.

0.2454
Finance score · higher is better

50.1
General score · higher is better

243.9 tps
Generation speed

₹144.85
Input price / 1M tokens

Why compare it: Speed comparison: the retained snapshot reports 243.9 output tokens per second. Useful when throughput matters.

View 1 configuration + evidence

Retained benchmark rate; exact SKU price needs claim-level verification

COHERE

Cohere Command A

Provider only

1 benchmark configuration

A Cohere model documented by the provider. This report does not have an exact matching row in the retained independent benchmark snapshot.

0.2850
Finance score · higher is better

48.5
General score · higher is better

112.0 tps
Generation speed

₹48.28
Input price / 1M tokens

> **Why compare it:** Finance comparison: its retained TAU Banking score is 0.2850. Treat this as a shortlist signal, not production proof.

View 1 configuration + evidence

Retained benchmark rate; exact SKU price needs claim-level verification

ANTHROPIC

Provider + benchmark

Claude Sonnet 4.6

1 benchmark configuration

An Anthropic model row. The provider documents this model family, and the retained independent benchmark snapshot contains the exact row name. The retained benchmark row uses the setting: Adaptive Reasoning, Max Effort.

0.3052
Finance score · higher is better

47.2
General score · higher is better

Not available
Generation speed

₹289.70
Input price / 1M tokens

> **Why compare it:** Finance comparison: its retained TAU Banking score is 0.3052. Treat this as a shortlist signal, not production proof.

View 1 configuration + evidence

Retained benchmark rate; exact SKU price needs claim-level verification

AMAZON

Provider only

Amazon Nova Pro

1 benchmark configuration

An Amazon model documented by the provider. This report does not have an exact matching row in the retained independent benchmark snapshot.

0.2710
Finance score · higher is better

47.2
General score · higher is better

145.0 tps
Generation speed

₹77.25
Input price / 1M tokens

Why compare it: Finance comparison: its retained TAU Banking score is 0.2710. Treat this as a shortlist signal, not production proof.

View 1 configuration + evidence

Retained benchmark rate; exact SKU price needs claim-level verification

GOOGLE

Provider + benchmark

Gemini 3.1 Pro Preview

1 benchmark configuration

A Google model row. The provider documents this model family, and the retained independent benchmark snapshot contains the exact row name.

0.1649
Finance score · higher is better

46.5
General score · higher is better

132.2 tps
Generation speed

₹193.14
Input price / 1M tokens

Why compare it: Use it as a cost and benchmark comparison row. No production recommendation is implied.

View 1 configuration + evidence

Retained benchmark rate; exact SKU price needs claim-level verification

ALIBABA

Provider + benchmark

Qwen3.7 Max

1 benchmark configuration

An Alibaba model row. The provider documents this model family, and the retained independent benchmark snapshot contains the exact row name.

0.1093
Finance score · higher is better

46.0
General score · higher is better

199.6 tps
Generation speed

₹241.42
Input price / 1M tokens

Why compare it: Speed comparison: the retained snapshot reports 199.6 output tokens per second. Useful when throughput matters.

View 1 configuration + evidence

Retained benchmark rate; exact SKU price needs claim-level verification

KIMI

Kimi K2.6

Provider + benchmark

1 benchmark configuration

A general-purpose Kimi model. The provider documents text, image and video input, thinking modes, tool calls and a 262,144-token context window.

0.2062
Finance score · higher is better

44.2
General score · higher is better

Not available
Generation speed

₹91.74
Input price / 1M tokens

> **Why compare it:** Useful as a lower-cost finance and agent-workflow candidate. Its retained TAU Banking score is 0.2062.

View 1 configuration + evidence

Verified exact official rate card

ALIBABA

Qwen3.7 Plus

Provider + benchmark

1 benchmark configuration

A lower-cost Qwen3.7 API tier listed by Alibaba Cloud Model Studio with the exact model ID qwen3.7-plus.

0.1787
Finance score · higher is better

39.0
General score · higher is better

52.8 tps
Generation speed

₹38.63
Input price / 1M tokens

> **Why compare it:** Useful when price matters more than using the Max tier. Its retained TAU Banking score is 0.1787.

View 1 configuration + evidence

Verified exact official rate card

GOOGLE

Gemini 3.5 Flash-Lite

Provider + benchmark

1 benchmark configuration

Google describes Gemini 3.5 Flash-Lite as its fastest and most cost-effective 3.5 model for high-throughput execution.

0.1649
Finance score · higher is better

36.5
General score · higher is better

362.2 tps
Generation speed

₹28.97
Input price / 1M tokens

> **Why compare it:** Useful for high-volume work. The retained benchmark snapshot reports 362.2 output tokens per second.

View 1 configuration + evidence

Verified exact official rate card

SECTION 8

Enterprise Compliance & EU AI Act Governance

EU AI Act Annex III, Article 15, SOC 2 Type II, HIPAA BAA and FedRAMP High are different governance mechanisms. This section separates legal scope, attestations, contracts, cloud authorization and SARVAX engineering policy.

Verified legal text

Annex III applicability is use-case specific

Annex III includes AI used to evaluate the creditworthiness or credit score of natural persons and AI used for risk assessment and pricing for natural persons in life and health insurance. A portfolio-analysis or wealth-advisory system is not automatically Annex III high-risk solely because it operates in finance; intended purpose and material influence must be assessed.

Official EU Annex III

Verified legal text

Article 15: accuracy, robustness and cybersecurity

Article 15 requires high-risk AI systems to achieve an appropriate level of accuracy, robustness and cybersecurity throughout their lifecycle, declare relevant accuracy metrics and be resilient to errors, faults and attacks. It does not name INT4, FP8 or BF16.

European Commission Article 15

▶ **SARVAX evidence required before EU deployment**

Third-party attestation

SOC 2 Type II

SOC 2 Type II concerns controls for a defined service-organization system over a review period. It is not a compliance badge inherited by an individual model. Review the report scope, Trust Services Criteria, examination period, exceptions and bridge coverage.

AICPA SOC resources

Contractual mechanism

HIPAA Business Associate Agreement

Where PHI is handled, a covered entity must obtain appropriate assurances through a business-associate contract and confirm the relevant service and data flow are in scope. "HIPAA eligible" model availability alone is insufficient.

HHS business-associate guidance

▶ Evidence required

AWS statement verified

Amazon Bedrock in GovCloud

AWS states that Amazon Bedrock is a FedRAMP High authorized service in AWS GovCloud (US-West). That statement applies to the authorized cloud service and boundary—not generically to every model, application architecture or customer deployment.

AWS Bedrock security & compliance

100% sovereignty claim rejected

Private VPC is not automatic air-gapping

Self-hosting a model in a VPC does not by itself guarantee air-gapping or "100% cloud data sovereignty." Sovereignty depends on region, identity, egress, private endpoints, encryption keys, logs, backups, telemetry, support access, subprocessors and the model-artifact supply chain.

AWS GovCloud

▶ FedRAMP High review boundary

Not statutory text

Article 15 does not prohibit INT4

The statement “INT4 is strictly prohibited and FP8/BF16 are required” is not present in Article 15. Quantization can affect numeric fidelity, but the legal requirement is outcome-based accuracy, robustness and cybersecurity.

[Verify Article 15]

SARVAX engineering policy

Material financial-risk scoring precision gate

SARVAX policy: do not deploy INT4 for material financial-risk scoring unless model-specific validation demonstrates no material degradation against the approved reference precision. FP8/BF16 may be approved only after the same tests; neither is automatically safe.

▶ Required precision-validation pack