

Evidence-scoped catalog of 24 model cards across 40 retained benchmark configurations. Proposed routing is evaluated separately and may include provider-verified models without catalog-grade benchmark matching.

Evidence:

FX BASIS

PROOF STATE

# Verdict: Conditional Hybrid Cascade — Evidence Gates Remain

A hybrid cascade remains a valid **architecture pattern**, not a verified model-selection decision. Twenty-one unique models have both provider-family documentation and at least one exact retained benchmark row; three need more proof. Financial savings, provider rate limits and production SLAs require exact SKU, contract and trace evidence before GTM approval.

**Decision:** architecture exploration may continue; procurement, financial commitment and regulated production routing remain conditional.

**Evidence boundary:** provider pages establish model-family documentation. Artificial Analysis metrics are from a retained API snapshot that could not be freshly refreshed because the API credential was unavailable. A benchmark row does not establish provider availability, SLA, compliance or customer access.

21

DUAL-PROOF MODEL  
CARDS

3

WEALTH-ADVISORY  
JOURNEYS

12

DOMAIN AGENT  
ROLES

3

PENDING  
VERIFICATION

# Executive Product Verdict: Wealth Advisory Architecture

A **Product and Enterprise Margins POV** evaluating the frontier landscape for long-running financial wealth workflows.

SARVAX WEALTH-ADVISORY ARCHITECTURE

12 agents · 3 advisory journeys

## Models underneath the product—not the product itself

Wealth advisory work arrives through meetings, CRM history, email, WhatsApp, holdings, market data, KYC records, tasks and client documents. SARVAX owns the journey: intake, agent orchestration, memory, approvals, integrations and writeback. Foundation models perform bounded routing, context assembly, vision and reasoning roles underneath it.

**Product recommendation:** Design around advisor journeys—not PDF types. Keep deterministic calculations, policy enforcement, authorization and human approval outside the model chain.

Relationship intake

Meeting Assistant + CRM/MCP context



The Router

Gemini 3.5 Flash-Lite



The Reader

DeepSeek V4 Pro



Relationship agents

Pre/Post Meeting + Client Strategist



### Advisor approval

Workflow 2.0 approval gate



### System writeback

CRM, tasks, notifications, artifacts

Turns meetings and ongoing client communication into advisor-ready context, follow-ups and controlled system updates.

#### SELECTED SARVAX JOURNEY

##### ADVISOR OUTCOME

Every advisor starts prepared, leaves with structured actions and keeps the client record current.

##### PRODUCT SURFACES

Meeting Assistant

Agent Builder

Agent Memory

Workflow 2.0

MCP / integrations layer

Notifications and writeback

## Client Relationship Intelligence

##### SARVAX AGENTS

Pre-Meeting Analysis Agent

Post-Meeting Analysis Agent

Client Intelligence Strategist

##### INPUTS ACROSS THE JOURNEY

Meeting Assistant transcripts

Email and WhatsApp context

CRM household history

Advisor notes and tasks

##### CONTROL BOUNDARY

Meeting intelligence drafts context and actions. The advisor approves material client communication and every financial recommendation.

BOUNDED MODEL ROLES UNDERNEATH THESE JOURNEYS

## Moonshot AI

**Product role:** Produces bounded portfolio, tax and compliance analysis drafts inside SARVAX agents. Deterministic calculators and advisor approval own the final decision.

**Why:** Its retained TAU Banking result is 0.3340, the highest value in this report's retained 40-configuration comparison set. This is a shortlist signal, not proof of production accuracy.

**Boundary:** Use only where a wealth-advisory journey needs deeper financial reasoning. Meeting notes, routing and routine extraction should not pay this cost or inherit this latency.

**0.3340**

TAU Banking  
Retained independent snapshot

**33.1 tps**

Speed  
Retained independent snapshot

**₹289.70**

Input / 1M  
Retained benchmark rate; exact SKU price  
needs claim-level verification

[Task testing required](#)[Provider proof](#) • [Benchmark](#)

## DeepSeek

**Product role:** Normalizes high-volume text from meeting transcripts, email, WhatsApp exports, CRM notes, research and text-native statements before SARVAX agents reason over it.

**Why:** The official rate card verifies the exact SKU, low cache-hit pricing, JSON output, tool calls and a 1M context window—useful for large client-context assembly.

**Boundary:** It is a text-context worker, not the product. Image-origin evidence still needs the Vision Gate; extracted facts must retain source references and pass journey-specific validation.

**0.2371–0.2577**

TAU Banking range  
Retained independent configurations

**31.2–44.3**

General score range  
Retained independent configurations

**₹42.01**

Cache-miss input / 1M  
Official DeepSeek rate card

[Official SKU and price verified; task quality unproven](#)[Provider proof](#) • [Benchmark](#)

Google

**Product role:** Identifies the wealth-advisory journey, urgency, household context and required SARVAX agent/workflow before expensive processing begins.

**Why:** Google describes it as its fastest, most cost-effective 3.5 model for high-throughput execution. The retained snapshot reports 362.2 tps; its verified input rate is ₹28.97/1M.

**Boundary:** Use for routing, classification and low-risk extraction checks—not investment, tax or compliance conclusions. Low-confidence routing must fall back to review.

|                                                            |                                                            |                                                               |
|------------------------------------------------------------|------------------------------------------------------------|---------------------------------------------------------------|
| <b>362.2 tps</b><br>Speed<br>Retained independent snapshot | <b>₹28.97</b><br>Input / 1M<br>Verified official rate card | <b>0.1649</b><br>TAU Banking<br>Retained independent snapshot |
|------------------------------------------------------------|------------------------------------------------------------|---------------------------------------------------------------|

Conditional gate; task testing required

[Provider proof](#) • [Benchmark](#)

Google

**Product role:** Handles image-origin evidence—scanned KYC records, photographed forms, charts and statement pages—when a broader wealth-advisory journey actually needs it.

**Why:** Google’s model page describes Gemini 3.6 Flash as supporting agentic and multimodal tasks. The retained snapshot reports 243.9 tps.

**Boundary:** Multimodal identity does not prove financial-document OCR accuracy. Validate tables, decimal placement, signatures and low-quality scans; escalate low-confidence evidence.

|                                                            |                                                               |                                                                                                                           |
|------------------------------------------------------------|---------------------------------------------------------------|---------------------------------------------------------------------------------------------------------------------------|
| <b>243.9 tps</b><br>Speed<br>Retained independent snapshot | <b>50.1</b><br>General score<br>Retained independent snapshot | <b>Exact SKU rate not established</b><br>Price<br>Retained benchmark rate; exact SKU price needs claim-level verification |
|------------------------------------------------------------|---------------------------------------------------------------|---------------------------------------------------------------------------------------------------------------------------|

Multimodal identity verified; OCR accuracy unproven

[Provider proof](#) • [Benchmark](#)

► Evidence boundaries and QA corrections

# SARVAX Wealth-Advisory Agent Operating Matrix

Twelve product agents mapped across three advisor journeys. No invented time-saved, accuracy, SLA or unit-economics claims are shown. Economics stay in the scenario calculator until measured traces exist.

CLIENT RELATIONSHIP INTELLIGENCE

Turns meetings and ongoing client communication into advisor-ready context, follow-ups and controlled system updates.

- Pre-Meeting Analysis Agent
- Post-Meeting Analysis Agent
- Client Intelligence Strategist

|                  |                                                                                                                                         |
|------------------|-----------------------------------------------------------------------------------------------------------------------------------------|
| PRODUCT SURFACES | Meeting Assistant · Agent Builder · Agent Memory · Workflow 2.0 · MCP / integrations layer · Notifications and writeback                |
| INPUTS           | Meeting Assistant transcripts · Email and WhatsApp context · CRM household history · Advisor notes and tasks                            |
| CONTROL BOUNDARY | Meeting intelligence drafts context and actions. The advisor approves material client communication and every financial recommendation. |

EVIDENCE STATUS

Frontend product contracts observed. Backend execution, connector availability, production reliability and branch freshness remain unproven.

PORTFOLIO & MARKET INTELLIGENCE

Combines holdings, market context and client objectives into review-ready portfolio and tax-analysis drafts.

- Portfolio Analyst
- Market Intelligence Agent

|                  |                                                                                                                                  |
|------------------|----------------------------------------------------------------------------------------------------------------------------------|
| PRODUCT SURFACES | MCP / integrations layer · OneChat projects and artifacts · Agent Builder · Skills · Workflow 2.0 · Approval and audit trail     |
| INPUTS           | Holdings and transactions · Market and research feeds · Risk profile and goals · Statements and tax-lot data                     |
| CONTROL BOUNDARY | Models can draft analysis. Deterministic portfolio/tax engines validate the math; the advisor owns suitability and final action. |

EVIDENCE STATUS

Frontend product contracts observed. Backend execution, connector availability, production reliability and branch freshness remain unproven.

COMPLIANCE & ADVISOR OPERATIONS

Routes KYC, review cadence, pipeline, workload and document events through bounded agents and approval-controlled workflows.

- Cadence Ops Lead
- Campaign Audience Architect
- Pipeline & Escalation Tracker
- Team Performance & Workload Insights Agent
- KYC Intelligence Officer
- KYC/CDD QA Auditor
- Document Intelligence & Auto-Population Agent

|                  |                                                                                                                                                            |
|------------------|------------------------------------------------------------------------------------------------------------------------------------------------------------|
| PRODUCT SURFACES | Agent Builder · Agent Memory · Workflow 2.0 parallel groups · Approval gates · Skills · MCP / integrations layer · Notifications                           |
| INPUTS           | KYC/CDD records · Client and account changes · Review cadence and pipeline events · Forms, messages and team tasks                                         |
| CONTROL BOUNDARY | Policy rules and deterministic validation remain outside the model chain. Compliance exceptions and material account changes require named human approval. |

EVIDENCE STATUS

Frontend product contracts observed. Backend execution, connector availability, production reliability and branch freshness remain unproven.



# Standard-Token Cost Scenario

Compares standard token rates across frontier models in the repository catalog using exact Decimal arithmetic and the official ECB USD/INR reference rate. Workload volumes are editable planning assumptions—not production invoices or observed SARVAX usage.

**Verified inputs**

Catalog rate cards and ECB 2026-07-24 USD/INR cross-rate.

**Scenario inputs**

Input tokens, output tokens and monthly execution volume.

**Excluded**

Cache, batch, retries, tools, storage, platform fees, taxes and negotiated contracts.

**Formula:**  $((\text{input tokens} \times \text{input USD/1M}) + (\text{output tokens} \times \text{output USD/1M})) \times \text{USD/INR}$ . Model-specific context and pricing tiers are enforced before a result is shown.

WORKLOAD SCENARIO PRESET

Client Relationship  
20k in · 2k out

Portfolio Review  
75k in · 8k out

Compliance Case  
50k in · 4k out

MODEL A (BASELINE)

Gemini 3.5 Flash-Lite — \$0.30/\$2.50 [Verified] ▼

↔ Swap

MODEL B (COMPARISON)

DeepSeek V4 Pro — \$0.435/\$0.87 [Verified] ▼

INPUT TOKENS PER RUN

75000tok

OUTPUT TOKENS PER RUN

8000tok

MONTHLY EXECUTION VOLUME

10000runs

FX basis: ECB 2026-07-24 — ₹96.567636/USD. Latest business-day official cross-rate.

Calculated with ECB 2026-07-24 FX (₹96.57/\$1) and verified catalog rates. Gemini 3.5 Flash-Lite: Global standard inference; official Google rate card. DeepSeek V4 Pro: Standard cache-miss input and output; official DeepSeek rate card.

CALCULATED STANDARD-TOKEN COST

₹ INR · ECB FX

Global standard inference; official Google rate card

PER RUN

₹4.10

PER 1,000 RUNS

₹4,104.12

MONTHLY EXPENDITURE

₹41,041.25 / mo

Standard cache-miss input and output; official DeepSeek rate card

PER RUN

₹3.82

PER 1,000 RUNS

₹3,822.63

MONTHLY EXPENDITURE

₹38,226.30 / mo

CALCULATED PRICE DIFFERENCE

DeepSeek V4 Pro is lower for these planning inputs. This is a rate-only difference—not measured savings.

ANNUALIZED DIFFERENCE

INTERPRETATION

## SECTION 4

# Wealth-Advisory Control-Flow Sandbox

Tests where product orchestration, deterministic gates and advisor approval must sit before a POC is approved. It demonstrates control order—not production execution quality.

Client Relationship Intelligence



Scenario progress

Ready

Evidence boundary

Stage sequence only · No measured tokens, cost, latency, cache, retries or SLA

DECISION UNDER TEST

# Can SARVAX preserve evidence, ownership and human decision rights across the complete meeting journey?

From permitted client context to advisor-approved follow-up and controlled system writeback.

PRODUCT OWNER

Meeting Assistant + Workflow 2.0

DETERMINISTIC OWNER

Policy, authorization and write-integrity services

HUMAN DECISION RIGHT

Named advisor

EXIT GATE

Release only after consent, tenant isolation, source lineage, approval and writeback controls pass representative trace tests.

RELEASE EVIDENCE STILL REQUIRED

Connector execution, persistence, retry behavior and production reliability are not yet proven with representative backend traces.

Control-flow chain

Governance and approval scenario — not production telemetry

Pending Active Passed

1 Permitted client-context intake PRODUCT PENDING

Calendar, CRM and approved communication context are collected before preparation.

**Required gate:** Required: consent, tenant access, source timestamps and retention policy.

2 Pre-Meeting Analysis Agent AGENT PENDING

Drafts a source-linked preparation brief and explicit data gaps.

**Required gate:** Required: unsupported client facts blocked; every material claim links to retained evidence.

3 Meeting capture and artifact finalization PRODUCT PENDING

Meeting Assistant creates the permitted transcript and meeting artifacts.

**Required gate:** Required: participant consent, speaker attribution review and artifact access policy.

4

**Post-Meeting Analysis Agent**

AGENT

PENDING

Drafts facts, decisions, risks and action items from the retained artifact.

**Required gate:** Required: transcript facts remain separate from model interpretation.

5

**Client Intelligence Strategist**

AGENT

PENDING

Drafts next-best-action options for advisor review.

**Required gate:** Required: no autonomous product recommendation, promise or client communication.

6

**Advisor decision gate**

HUMAN

PENDING

The named advisor accepts, edits or rejects facts and proposed actions.

**Required gate:** Required before any client communication or material CRM/task update.

7

**Controlled Workflow 2.0 writeback**

TOOL


PENDING

Only approved records are submitted to authorized connected systems.


**Required gate:** Release evidence required: idempotency key, immutable audit event, retry/dead-letter policy and compensating action.

# Four Evidence-Backed Comparison Charts


Charts read the central runtime dataset. Each chart discloses population, metric direction, snapshot boundary and unavailable data. Accessible tables mirror every plotted value.



Highest retained TAU Banking configuration per model; higher is better; retained snapshot.




Counts unique model cards by provider/benchmark proof type.





Exact public standard-token rates converted with the retained ECB FX. Qwen uses the  $\leq 256K$  input tier in this chart; the calculator applies its higher verified tier above 256K.



## SECTION 6

# Official Public Proofs & Evidence Boundaries

33 usable source snapshots are content-addressed. Model-family identity comes from official providers; benchmark values come from a retained Artificial Analysis API snapshot and are labelled accordingly.

### ▼ Official provider documentation — current retrieval

Anthropic, OpenAI, Google, DeepSeek, Moonshot/Kimi, xAI, Z.ai, Alibaba Cloud, Cohere and AWS pages were fetched, redirected URLs retained, and SHA-256 hashes recorded. Click any model proof badge in Section 7 to inspect its exact provider link.

### ► Independent benchmark snapshot — retained, not freshly refreshed

### ► Pricing and FX boundary

# 24 Unique Model Cards — 40 Benchmark Configurations

Each model appears once. Adaptive Reasoning and effort labels are mapped inside the card as benchmark configurations, not counted as separate models. Every card explains the score range, evidence and open gaps in simple words. INR values use the ECB 2026-07-24 cross-rate of ₹96.567636/USD.

ANTHROPIC

Provider + benchmark

## Claude Opus 5

5 benchmark configurations

One Anthropic model represented by 5 retained benchmark configurations. The configuration labels are grouped here instead of being counted as separate model cards.

0.2330–0.3278

Finance score range · higher is better

50.6–60.7

General score range · higher is better

43.9–63.8 tps

Generation speed range

₹482.84

Input price / 1M tokens

**Why compare it:** Use the ranges to see how benchmark settings changed the result. Open the guide to compare all 5 configurations side by side.

View 5 configurations + evidence

Retained benchmark rate; exact SKU price needs claim-level verification

ANTHROPIC

Provider + benchmark

## Claude Fable 5

1 benchmark configuration

An Anthropic model row. The provider documents this model family, and the retained independent benchmark snapshot contains the exact row name. The retained benchmark row uses the setting: Adaptive Reasoning, Max Effort, Opus 4.8 Fallback.

0.2680

Finance score · higher is better

59.9

General score · higher is better

58.3 tps

Generation speed

₹965.68

Input price / 1M tokens

**Why compare it:** Finance comparison: its retained TAU Banking score is 0.2680. Treat this as a shortlist signal, not production proof.

[View 1 configuration + evidence](#)

Retained benchmark rate; exact SKU price needs claim-level verification

OPENAI

[Provider + benchmark](#)

## GPT-5.6 Sol

5 benchmark configurations

One OpenAI model represented by 5 retained benchmark configurations. The configuration labels are grouped here instead of being counted as separate model cards.

0.2440–0.3299

Finance score range · higher is better

49.4–58.9

General score range · higher is better

60.0–73.9 tps

Generation speed range

₹482.84

Input price / 1M tokens

**Why compare it:** Use the ranges to see how benchmark settings changed the result. Open the guide to compare all 5 configurations side by side.

[View 5 configurations + evidence](#)

Retained benchmark rate; exact SKU price needs claim-level verification

KIMI

[Provider + benchmark](#)

## Kimi K3

1 benchmark configuration

A Kimi model row. The provider documents this model family, and the retained independent benchmark snapshot contains the exact row name.

0.3340

Finance score · higher is better

57.1

General score · higher is better

33.1 tps

Generation speed

₹289.70

Input price / 1M tokens

**Why compare it:** Finance comparison: its retained TAU Banking score is 0.3340. Treat this as a shortlist signal, not production proof.

[View 1 configuration + evidence](#)

Retained benchmark rate; exact SKU price needs claim-level verification

ANTHROPIC

[Provider + benchmark](#)

## Claude Opus 4.8

1 benchmark configuration

An Anthropic model row. The provider documents this model family, and the retained independent benchmark snapshot contains the exact row name. The retained benchmark row uses the setting: Adaptive Reasoning, Max Effort.

0.2763

Finance score · higher is better

55.7

General score · higher is better

62.5 tps

Generation speed

₹482.84

Input price / 1M tokens

**Why compare it:** Finance comparison: its retained TAU Banking score is 0.2763. Treat this as a shortlist signal, not production proof.

[View 1 configuration + evidence](#)

Retained benchmark rate; exact SKU price needs claim-level verification

OPENAI

[Provider + benchmark](#)

## GPT-5.6 Terra

4 benchmark configurations

One OpenAI model represented by 4 retained benchmark configurations. The configuration labels are grouped here instead of being counted as separate model cards.

0.1938–0.3175

Finance score range · higher is better

45.6–55.0

General score range · higher is better

114.6–128.0 tps

Generation speed range

₹241.42

Input price / 1M tokens

**Why compare it:** Use the ranges to see how benchmark settings changed the result. Open the guide to compare all 4 configurations side by side.

[View 4 configurations + evidence](#)

Retained benchmark rate; exact SKU price needs claim-level verification

OPENAI

[Provider + benchmark](#)

## GPT-5.5

3 benchmark configurations

One OpenAI model represented by 3 retained benchmark configurations. The configuration labels are grouped here instead of being counted as separate model cards.

**0.2577–0.3134**

Finance score range · higher is better

**50.4–54.8**

General score range · higher is better

**Not available**

Generation speed range

**₹482.84**

Input price / 1M tokens

**Why compare it:** Use the ranges to see how benchmark settings changed the result. Open the guide to compare all 3 configurations side by side.

[View 3 configurations + evidence](#)

Retained benchmark rate; exact SKU price needs claim-level verification

SPACEXAI

[Provider + benchmark](#)

## Grok 4.5

1 benchmark configuration

A SpaceXAI model row. The provider documents this model family, and the retained independent benchmark snapshot contains the exact row name. The retained benchmark row uses the setting: high.

**0.3258**

Finance score · higher is better

**53.8**

General score · higher is better

**55.8 tps**

Generation speed

**₹193.14**

Input price / 1M tokens

**Why compare it:** Finance comparison: its retained TAU Banking score is 0.3258. Treat this as a shortlist signal, not production proof.

[View 1 configuration + evidence](#)

ANTHROPIC

Provider + benchmark

# Claude Opus 4.7

1 benchmark configuration

An Anthropic model row. The provider documents this model family, and the retained independent benchmark snapshot contains the exact row name. The retained benchmark row uses the setting: Adaptive Reasoning, Max Effort.

0.2887

Finance score · higher is better

53.5

General score · higher is better

Not available

Generation speed

₹482.84

Input price / 1M tokens

**Why compare it:** Finance comparison: its retained TAU Banking score is 0.2887. Treat this as a shortlist signal, not production proof.

View 1 configuration + evidence

ANTHROPIC

Provider + benchmark

# Claude Sonnet 5

1 benchmark configuration

An Anthropic model row. The provider documents this model family, and the retained independent benchmark snapshot contains the exact row name. The retained benchmark row uses the setting: Adaptive Reasoning, Max Effort.

0.2825

Finance score · higher is better

53.4

General score · higher is better

83.4 tps

Generation speed

₹193.14

Input price / 1M tokens

**Why compare it:** Finance comparison: its retained TAU Banking score is 0.2825. Treat this as a shortlist signal, not production proof.

View 1 configuration + evidence

OPENAI

[Provider + benchmark](#)

## GPT-5.4

1 benchmark configuration

An OpenAI model row. The provider documents this model family, and the retained independent benchmark snapshot contains the exact row name. The retained benchmark row uses the setting: xhigh.

**0.3031**

Finance score · higher is better

**51.4**

General score · higher is better

**Not available**

Generation speed

**₹241.42**

Input price / 1M tokens

**Why compare it:** Finance comparison: its retained TAU Banking score is 0.3031. Treat this as a shortlist signal, not production proof.

[View 1 configuration + evidence](#)

Retained benchmark rate; exact SKU price needs claim-level verification

OPENAI

[Provider + benchmark](#)

## GPT-5.6 Luna

3 benchmark configurations

One OpenAI model represented by 3 retained benchmark configurations. The configuration labels are grouped here instead of being counted as separate model cards.

**0.2227–0.2722**

Finance score range · higher is better

**46.1–51.2**

General score range · higher is better

**161.2–174.3 tps**

Generation speed range

**₹96.57**

Input price / 1M tokens

**Why compare it:** Use the ranges to see how benchmark settings changed the result. Open the guide to compare all 3 configurations side by side.

[View 3 configurations + evidence](#)

Retained benchmark rate; exact SKU price needs claim-level verification

Z AI

[Provider + benchmark](#)

## GLM-5.2

1 benchmark configuration



A Z AI model row. The provider documents this model family, and the retained independent benchmark snapshot contains the exact row name. The retained benchmark row uses the setting: max.

0.2680

Finance score · higher is better

51.1

General score · higher is better

156.7 tps

Generation speed

₹135.19

Input price / 1M tokens

**Why compare it:** Finance comparison: its retained TAU Banking score is 0.2680. Treat this as a shortlist signal, not production proof.

[View 1 configuration + evidence](#)

Retained benchmark rate; exact SKU price needs claim-level verification

META

Benchmark only

## Muse Spark 1.1

1 benchmark configuration

A model name found in the retained independent benchmark snapshot, but current official provider documentation was not established. The retained benchmark row uses the setting: xhigh.

0.2515

Finance score · higher is better

50.6

General score · higher is better

123.9 tps

Generation speed

₹120.71

Input price / 1M tokens

**Why compare it:** Finance comparison: its retained TAU Banking score is 0.2515. Treat this as a shortlist signal, not production proof.

[View 1 configuration + evidence](#)

Retained benchmark rate; exact SKU price needs claim-level verification

GOOGLE

Provider + benchmark

## Gemini 3.5 Flash

2 benchmark configurations

One Google model represented by 2 retained benchmark configurations. The configuration labels are grouped here instead of being counted as separate model cards.

0.2536

Finance score range · higher is better

45.4–50.2

General score range · higher is better

250.3–265.2 tps

Generation speed range

₹144.85

Input price / 1M tokens

**Why compare it:** Use the ranges to see how benchmark settings changed the result. Open the guide to compare all 2 configurations side by side.

[View 2 configurations + evidence](#)

Retained benchmark rate; exact SKU price needs claim-level verification

GOOGLE

[Provider + benchmark](#)

## Gemini 3.6 Flash

1 benchmark configuration

A Google model row. The provider documents this model family, and the retained independent benchmark snapshot contains the exact row name. The retained benchmark row uses the setting: high.

0.2454

Finance score · higher is better

50.1

General score · higher is better

243.9 tps

Generation speed

₹144.85

Input price / 1M tokens

**Why compare it:** Speed comparison: the retained snapshot reports 243.9 output tokens per second. Useful when throughput matters.

[View 1 configuration + evidence](#)

Retained benchmark rate; exact SKU price needs claim-level verification

COHERE

[Provider only](#)

## Cohere Command A

1 benchmark configuration

A Cohere model documented by the provider. This report does not have an exact matching row in the retained independent benchmark snapshot.

0.2850

Finance score · higher is better

48.5

General score · higher is better

112.0 tps

Generation speed

₹48.28

Input price / 1M tokens

**Why compare it:** Finance comparison: its retained TAU Banking score is 0.2850. Treat this as a shortlist signal, not production proof.

[View 1 configuration + evidence](#)

Retained benchmark rate; exact SKU price needs claim-level verification

ANTHROPIC

[Provider + benchmark](#)

## Claude Sonnet 4.6

1 benchmark configuration

An Anthropic model row. The provider documents this model family, and the retained independent benchmark snapshot contains the exact row name. The retained benchmark row uses the setting: Adaptive Reasoning, Max Effort.

0.3052

Finance score · higher is better

47.2

General score · higher is better

Not available

Generation speed

₹289.70

Input price / 1M tokens

**Why compare it:** Finance comparison: its retained TAU Banking score is 0.3052. Treat this as a shortlist signal, not production proof.

[View 1 configuration + evidence](#)

Retained benchmark rate; exact SKU price needs claim-level verification

AMAZON

[Provider only](#)

## Amazon Nova Pro

1 benchmark configuration

An Amazon model documented by the provider. This report does not have an exact matching row in the retained independent benchmark snapshot.

0.2710

Finance score · higher is better

47.2

General score · higher is better

145.0 tps

Generation speed

₹77.25

Input price / 1M tokens

**Why compare it:** Finance comparison: its retained TAU Banking score is 0.2710. Treat this as a shortlist signal, not production proof.

[View 1 configuration + evidence](#)

Retained benchmark rate; exact SKU price needs claim-level verification

GOOGLE

[Provider + benchmark](#)

## Gemini 3.1 Pro Preview

1 benchmark configuration

A Google model row. The provider documents this model family, and the retained independent benchmark snapshot contains the exact row name.

**0.1649**

Finance score · higher is better

**46.5**

General score · higher is better

**132.2 tps**

Generation speed

**₹193.14**

Input price / 1M tokens

**Why compare it:** Use it as a cost and benchmark comparison row. No production recommendation is implied.

[View 1 configuration + evidence](#)

Retained benchmark rate; exact SKU price needs claim-level verification

ALIBABA

[Provider + benchmark](#)

## Qwen3.7 Max

1 benchmark configuration

An Alibaba model row. The provider documents this model family, and the retained independent benchmark snapshot contains the exact row name.

**0.1093**

Finance score · higher is better

**46.0**

General score · higher is better

**199.6 tps**

Generation speed

**₹241.42**

Input price / 1M tokens

**Why compare it:** Speed comparison: the retained snapshot reports 199.6 output tokens per second. Useful when throughput matters.

[View 1 configuration + evidence](#)

KIMI

[Provider + benchmark](#)

## Kimi K2.6

1 benchmark configuration

A general-purpose Kimi model. The provider documents text, image and video input, thinking modes, tool calls and a 262,144-token context window.

**0.2062**

Finance score · higher is better

**44.2**

General score · higher is better

**Not available**

Generation speed

**₹91.74**

Input price / 1M tokens

**Why compare it:** Useful as a lower-cost finance and agent-workflow candidate. Its retained TAU Banking score is 0.2062.

[View 1 configuration + evidence](#)

Verified exact official rate card

ALIBABA

[Provider + benchmark](#)

## Qwen3.7 Plus

1 benchmark configuration

A lower-cost Qwen3.7 API tier listed by Alibaba Cloud Model Studio with the exact model ID qwen3.7-plus.

**0.1787**

Finance score · higher is better

**39.0**

General score · higher is better

**52.8 tps**

Generation speed

**₹38.63**

Input price / 1M tokens

**Why compare it:** Useful when price matters more than using the Max tier. Its retained TAU Banking score is 0.1787.

[View 1 configuration + evidence](#)

Verified exact official rate card

GOOGLE

[Provider + benchmark](#)

## Gemini 3.5 Flash-Lite

1 benchmark configuration

Google describes Gemini 3.5 Flash-Lite as its fastest and most cost-effective 3.5 model for high-throughput execution.

0.1649

Finance score · higher is better

36.5

General score · higher is better

362.2 tps

Generation speed

₹28.97

Input price / 1M tokens

**Why compare it:** Useful for high-volume work. The retained benchmark snapshot reports 362.2 output tokens per second.

View 1 configuration + evidence

Verified exact official rate card

# Enterprise Compliance & EU AI Act Governance

EU AI Act Annex III, Article 15, SOC 2 Type II, HIPAA BAA and FedRAMP High are different governance mechanisms. This section separates legal scope, attestations, contracts, cloud authorization and SARVAX engineering policy.

Verified legal text

## Annex III applicability is use-case specific

Annex III includes AI used to evaluate the creditworthiness or credit score of natural persons and AI used for risk assessment and pricing for natural persons in life and health insurance. A portfolio-analysis or wealth-advisory system is not automatically Annex III high-risk solely because it operates in finance; intended purpose and material influence must be assessed.

[Official EU Annex III](#)

Verified legal text

## Article 15: accuracy, robustness and cybersecurity

Article 15 requires high-risk AI systems to achieve an appropriate level of accuracy, robustness and cybersecurity throughout their lifecycle, declare relevant accuracy metrics and be resilient to errors, faults and attacks. It does not name INT4, FP8 or BF16.

[European Commission Article 15](#)

► SARVAX evidence required before EU deployment

Third-party attestation

SOC 2 Type II

SOC 2 Type II concerns controls for a defined service-organization system over a review period. It is not a compliance badge inherited by an individual model. Review the report scope, Trust Services Criteria, examination period, exceptions and bridge coverage.

[AICPA SOC resources](#)

Contractual mechanism

## HIPAA Business Associate Agreement

Where PHI is handled, a covered entity must obtain appropriate assurances through a business-associate contract and confirm the relevant service and data flow are in scope. "HIPAA eligible" model availability alone is insufficient.

[HHS business-associate guidance](#)

### ► Evidence required

AWS statement verified

## Amazon Bedrock in GovCloud

AWS states that Amazon Bedrock is a FedRAMP High authorized service in AWS GovCloud (US-West). That statement applies to the authorized cloud service and boundary—not generically to every model, application architecture or customer deployment.

[AWS Bedrock security & compliance](#)

100% sovereignty claim rejected

## Private VPC is not automatic air-gapping

Self-hosting a model in a VPC does not by itself guarantee air-gapping or "100% cloud data sovereignty." Sovereignty depends on region, identity, egress, private endpoints, encryption keys, logs, backups, telemetry, support access, subprocessors and the model-artifact supply chain.

[AWS GovCloud](#)



► FedRAMP High review boundary

Not statutory text

## Article 15 does not prohibit INT4

The statement “INT4 is strictly prohibited and FP8/BF16 are required” is not present in Article 15. Quantization can affect numeric fidelity, but the legal requirement is outcome-based accuracy, robustness and cybersecurity.

[Verify Article 15](#)

SARVAX engineering policy

## Material financial-risk scoring precision gate

SARVAX policy: do not deploy INT4 for material financial-risk scoring unless model-specific validation demonstrates no material degradation against the approved reference precision. FP8/BF16 may be approved only after the same tests; neither is automatically safe.

► Required precision-validation pack